
PySpark API

Release 1.00

Wenqiang Feng

March 19, 2019

CONTENTS

1	Preface	3
1.1	About	3
1.1.1	About this API	3
1.2	Feedback and suggestions	3
2	Stat API	5
3	Regression API	13
4	Classification API	39
5	Clustering API	67
6	Recommendation API	85
7	Pipeline API	93
8	Tuning API	97
9	Evaluation API	103
10	Main Reference	109
	Bibliography	111
	Python Module Index	113
	Index	115



Welcome to my **PySpark API**! The PDF version can be downloaded from [HERE](#).

PREFACE

Chinese proverb

Good tools are prerequisite to the successful execution of a job. – old Chinese proverb

1.1 About

1.1.1 About this API

This document is the API for [PySpark]. All CopyRights are belongs to Databricks and Spark team. This document is generated automatically by using [sphinx](#).

1.2 Feedback and suggestions

Your comments and suggestions are highly appreciated. I am more than happy to receive corrections, suggestions or feedbacks through email (Wenqiang Feng: von198@gmail.com) for improvements.

STAT API

```
class pyspark.ml.stat.ChiSquareTest
```

Note: Experimental

Conduct Pearson's independence test for every feature against the label. For each feature, the (feature, label) pairs are converted into a contingency matrix for which the Chi-squared statistic is computed. All label and feature values must be categorical.

The null hypothesis is that the occurrence of the outcomes is statistically independent.

New in version 2.2.0.

static test (*dataset*, *featuresCol*, *labelCol*)

Perform a Pearson's independence test using dataset.

Parameters

- **dataset** – DataFrame of categorical labels and categorical features. Real-valued features will be treated as categorical for each distinct value.
- **featuresCol** – Name of features column in dataset, of type *Vector* (*VectorUDT*).
- **labelCol** – Name of label column in dataset, of any numerical type.

Returns DataFrame containing the test result for every feature against the label. This DataFrame will contain a single Row with the following fields:
- *pValues* : *Vector* - *degreesOfFreedom* : *Array[Int]* - *statistics* : *Vector* Each of these fields has one value per feature.

```
>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.stat import ChiSquareTest
>>> dataset = [[0, Vectors.dense([0, 0, 1])],
```

(continues on next page)

(continued from previous page)

```

...         [0, Vectors.dense([1, 0, 1])],
...         [1, Vectors.dense([2, 1, 1])],
...         [1, Vectors.dense([3, 1, 1])]]
>>> dataset = spark.createDataFrame(dataset, ["label",
↪ "features"])
>>> chiSqResult = ChiSquareTest.test(dataset, 'features',
↪ 'label')
>>> chiSqResult.select("degreesOfFreedom").collect()[0]
Row(degreesOfFreedom=[3, 1, 0])

```

New in version 2.2.0.

class pyspark.ml.stat.**Correlation****Note:** Experimental

Compute the correlation matrix for the input dataset of Vectors using the specified method. Methods currently supported: *pearson* (default), *spearman*.

Note: For Spearman, a rank correlation, we need to create an RDD[Double] for each column and sort it in order to retrieve the ranks and then join the columns back into an RDD[Vector], which is fairly costly. Cache the input Dataset before calling corr with *method = 'spearman'* to avoid recomputing the common lineage.

New in version 2.2.0.

static corr (dataset, column, method='pearson')

Compute the correlation matrix with specified method using dataset.

Parameters

- **dataset** – A Dataset or a DataFrame.
- **column** – The name of the column of vectors for which the correlation coefficient needs to be computed. This must be a column of the dataset, and it must contain Vector objects.
- **method** – String specifying the method to use for computing correlation. Supported: *pearson* (default), *spearman*.

Returns A DataFrame that contains the correlation matrix of the column of vectors. This DataFrame contains a single row and a single column of name '\$METHODNAME(\$COLUMN)'.

```

>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.stat import Correlation
>>> dataset = [[Vectors.dense([1, 0, 0, -2])],
...            [Vectors.dense([4, 5, 0, 3])],
...            [Vectors.dense([6, 7, 0, 8])],
...            [Vectors.dense([9, 0, 0, 1])]]
>>> dataset = spark.createDataFrame(dataset, ['features'])
>>> pearsonCorr = Correlation.corr(dataset, 'features',
↳ 'pearson').collect()[0][0]
>>> print(str(pearsonCorr).replace('nan', 'NaN'))
DenseMatrix([[ 1.          ,  0.0556...,      NaN,  0.4004...
↳ ],
              [ 0.0556...,  1.          ,      NaN,  0.9135...
↳ ],
              [      NaN,      NaN,  1.          ,
↳ NaN],
              [ 0.4004...,  0.9135...,      NaN,  1.          ]
↳ ])
>>> spearmanCorr = Correlation.corr(dataset, 'features',
↳ method='spearman').collect()[0][0]
>>> print(str(spearmanCorr).replace('nan', 'NaN'))
DenseMatrix([[ 1.          ,  0.1054...,      NaN,  0.4
↳ ],
              [ 0.1054...,  1.          ,      NaN,  0.9486...
↳ ],
              [      NaN,      NaN,  1.          ,
↳ NaN],
              [ 0.4          ,  0.9486...,      NaN,  1.          ]
↳ ])

```

New in version 2.2.0.

class pyspark.ml.stat.KolmogorovSmirnovTest

Note: Experimental

Conduct the two-sided Kolmogorov Smirnov (KS) test for data sampled from a continuous distribution.

By comparing the largest difference between the empirical cumulative distribution of the sample data and the theoretical distribution we can provide a test for the the null hypothesis that the sample data comes from that theoretical distribution.

New in version 2.4.0.

static test (*dataset, sampleCol, distName, *params*)

Conduct a one-sample, two-sided Kolmogorov-Smirnov test for probability distribution equality. Currently supports the normal distribution, taking as parameters the mean and standard deviation.

Parameters

- **dataset** – a Dataset or a DataFrame containing the sample of data to test.
- **sampleCol** – Name of sample column in dataset, of any numerical type.
- **distName** – a *string* name for a theoretical distribution, currently only support “norm”.
- **params** – a list of *Double* values specifying the parameters to be used for the theoretical distribution. For “norm” distribution, the parameters includes mean and variance.

Returns A DataFrame that contains the Kolmogorov-Smirnov test result for the input sampled data. This DataFrame will contain a single Row with the following fields: - *pValue* : *Double* - *statistic* : *Double*

```
>>> from pyspark.ml.stat import KolmogorovSmirnovTest
>>> dataset = [[-1.0], [0.0], [1.0]]
>>> dataset = spark.createDataFrame(dataset, ['sample'])
>>> ksResult = KolmogorovSmirnovTest.test(dataset, 'sample',
↪ 'norm', 0.0, 1.0).first()
>>> round(ksResult.pValue, 3)
1.0
>>> round(ksResult.statistic, 3)
0.175
>>> dataset = [[2.0], [3.0], [4.0]]
>>> dataset = spark.createDataFrame(dataset, ['sample'])
>>> ksResult = KolmogorovSmirnovTest.test(dataset, 'sample',
↪ 'norm', 3.0, 1.0).first()
>>> round(ksResult.pValue, 3)
1.0
>>> round(ksResult.statistic, 3)
0.175
```

New in version 2.4.0.

class pyspark.ml.stat.**Summarizer**

Note: Experimental

Tools for vectorized statistics on MLlib Vectors. The methods in this package provide various statistics for Vectors contained inside DataFrames. This class lets users pick the statistics they would like to extract for a given column.

```
>>> from pyspark.ml.stat import Summarizer
>>> from pyspark.sql import Row
>>> from pyspark.ml.linalg import Vectors
>>> summarizer = Summarizer.metrics("mean", "count")
>>> df = sc.parallelize([Row(weight=1.0, features=Vectors.dense(1.
    ↪0, 1.0, 1.0)),
    ...                  Row(weight=0.0, features=Vectors.dense(1.
    ↪0, 2.0, 3.0))]).toDF()
>>> df.select(summarizer.summary(df.features, df.weight)).
    ↪show(truncate=False)
+-----+
|aggregate_metrics(features, weight)|
+-----+
|[[1.0,1.0,1.0], 1]|
+-----+
<BLANKLINE>
>>> df.select(summarizer.summary(df.features)).show(truncate=False)
+-----+
|aggregate_metrics(features, 1.0)|
+-----+
|[[1.0,1.5,2.0], 2]|
+-----+
<BLANKLINE>
>>> df.select(Summarizer.mean(df.features, df.weight)).
    ↪show(truncate=False)
+-----+
|mean(features)|
+-----+
|[1.0,1.0,1.0]|
+-----+
<BLANKLINE>
>>> df.select(Summarizer.mean(df.features)).show(truncate=False)
+-----+
|mean(features)|
+-----+
|[1.0,1.5,2.0]|
+-----+
<BLANKLINE>
```

New in version 2.4.0.

static count (*col*, *weightCol=None*)
 return a column of count summary

New in version 2.4.0.

static max (*col*, *weightCol=None*)
return a column of max summary

New in version 2.4.0.

static mean (*col*, *weightCol=None*)
return a column of mean summary

New in version 2.4.0.

static metrics (**metrics*)

Given a list of metrics, provides a builder that it turns computes metrics from a column.

See the documentation of `[[Summarizer]]` for an example.

The following metrics are accepted (case sensitive):

- mean: a vector that contains the coefficient-wise mean.
- variance: a vector tha contains the coefficient-wise variance.
- count: the count of all vectors seen.
- numNonzeros: a vector with the number of non-zeros for each coefficients
- max: the maximum for each coefficient.
- min: the minimum for each coefficient.
- normL2: the Euclidian norm for each coefficient.
- normL1: the L1 norm of each coefficient (sum of the absolute values).

Parameters metrics – metrics that can be provided.

Returns an object of `pyspark.ml.stat.SummaryBuilder`

Note: Currently, the performance of this interface is about 2x~3x slower then using the RDD interface.

New in version 2.4.0.

static min (*col*, *weightCol=None*)
return a column of min summary

New in version 2.4.0.

static normL1 (*col*, *weightCol=None*)
return a column of normL1 summary

New in version 2.4.0.

static normL2 (*col*, *weightCol=None*)
return a column of normL2 summary

New in version 2.4.0.

static numNonZeros (*col*, *weightCol=None*)
return a column of numNonZero summary

New in version 2.4.0.

static variance (*col*, *weightCol=None*)
return a column of variance summary

New in version 2.4.0.

class pyspark.ml.stat.**SummaryBuilder** (*jSummaryBuilder*)

Note: Experimental

A builder object that provides summary statistics about a given column.

Users should not directly create such builders, but instead use one of the methods in `pyspark.ml.stat.Summarizer`

New in version 2.4.0.

summary (*featuresCol*, *weightCol=None*)

Returns an aggregate object that contains the summary of the column with the requested metrics.

Parameters

- **featuresCol** – a column that contains features Vector object.
- **weightCol** – a column that contains weight value. Default weight is 1.0.

Returns an aggregate column that contains the statistics. The exact content of this structure is determined during the creation of the builder.

New in version 2.4.0.

REGRESSION API

```
class pyspark.ml.regression.AFTSurvivalRegression (featuresCol='features',  
                                                    label-  
                                                    Col='label',  
                                                    prediction-  
                                                    Col='prediction',  
                                                    fitInter-  
                                                    cept=True,  
                                                    maxIter=100,  
                                                    tol=1e-06,  
                                                    censor-  
                                                    Col='censor',  
                                                    quantileProba-  
                                                    bilities=[0.01,  
                                                    0.05,      0.1,  
                                                    0.25,      0.5,  
                                                    0.75, 0.9, 0.95,  
                                                    0.99],    quan-  
                                                    tilesCol=None,  
                                                    aggregation-  
                                                    Depth=2)
```

Note: Experimental

Accelerated Failure Time (AFT) Model Survival Regression

Fit a parametric AFT survival regression model based on the Weibull distribution of the survival time.

See also:

[AFT Model](#)

```

>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0), 1.0),
...     (1e-40, Vectors.sparse(1, [], []), 0.0)], ["label",
→ "features", "censor"])
>>> aftsr = AFTSurvivalRegression()
>>> model = aftsr.fit(df)
>>> model.predict(Vectors.dense(6.3))
1.0
>>> model.predictQuantiles(Vectors.dense(6.3))
DenseVector([0.0101, 0.0513, 0.1054, 0.2877, 0.6931, 1.3863, 2.
→ 3026, 2.9957, 4.6052])
>>> model.transform(df).show()
+-----+-----+-----+-----+
|  label| features|censor|prediction|
+-----+-----+-----+-----+
|    1.0|    [1.0]|    1.0|        1.0|
|1.0E-40|(1, [], [])|    0.0|        1.0|
+-----+-----+-----+-----+
...
>>> aftsr_path = temp_path + "/aftsr"
>>> aftsr.save(aftsr_path)
>>> aftsr2 = AFTSurvivalRegression.load(aftsr_path)
>>> aftsr2.getMaxIter()
100
>>> model_path = temp_path + "/aftsr_model"
>>> model.save(model_path)
>>> model2 = AFTSurvivalRegressionModel.load(model_path)
>>> model.coefficients == model2.coefficients
True
>>> model.intercept == model2.intercept
True
>>> model.scale == model2.scale
True

```

New in version 1.6.0.

getCensorCol()

Gets the value of censorCol or its default value.

New in version 1.6.0.

getQuantileProbabilities()

Gets the value of quantileProbabilities or its default value.

New in version 1.6.0.

getQuantilesCol()

Gets the value of quantilesCol or its default value.

New in version 1.6.0.

setCensorCol (*value*)

Sets the value of `censorCol`.

New in version 1.6.0.

setParams (*featuresCol='features', labelCol='label', predictionCol='prediction', fitIntercept=True, maxIter=100, tol=1e-06, censorCol='censor', quantileProbabilities=[0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95, 0.99], quantilesCol=None, aggregationDepth=2*)
 setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", fitIntercept=True, maxIter=100, tol=1E-6, censorCol="censor", quantileProbabilities=[0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95, 0.99], quantilesCol=None, aggregationDepth=2):

New in version 1.6.0.

setQuantileProbabilities (*value*)

Sets the value of `quantileProbabilities`.

New in version 1.6.0.

setQuantilesCol (*value*)

Sets the value of `quantilesCol`.

New in version 1.6.0.

class pyspark.ml.regression.**AFTSurvivalRegressionModel** (*java_model=None*)

Note: Experimental

Model fitted by *AFTSurvivalRegression*.

New in version 1.6.0.

coefficients

Model coefficients.

New in version 2.0.0.

intercept

Model intercept.

New in version 1.6.0.

predict (*features*)

Predicted value

New in version 2.0.0.

predictQuantiles (*features*)
Predicted Quantiles

New in version 2.0.0.

scale
Model scale parameter.
New in version 1.6.0.

class pyspark.ml.regression.**DecisionTreeRegressor** (*featuresCol='features', labelCol='label', predictionCol='prediction', maxDepth=5, maxBins=32, minInstancesPerNode=1, minInfoGain=0.0, maxMemoryInMB=256, cacheNodeIds=False, checkpointInterval=10, impurity='variance', seed=None, varianceCol=None*)

Decision tree learning algorithm for regression. It supports both continuous and categorical features.

```
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> dt = DecisionTreeRegressor(maxDepth=2, varianceCol="variance")
>>> model = dt.fit(df)
>>> model.depth
1
>>> model.numNodes
3
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> model.numFeatures
```

(continues on next page)

(continued from previous page)

```

1
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
  ↪ "features"])
>>> model.transform(test0).head().prediction
0.0
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
  ↪)], ["features"])
>>> model.transform(test1).head().prediction
1.0
>>> dtr_path = temp_path + "/dtr"
>>> dt.save(dtr_path)
>>> dt2 = DecisionTreeRegressor.load(dtr_path)
>>> dt2.getMaxDepth()
2
>>> model_path = temp_path + "/dtr_model"
>>> model.save(model_path)
>>> model2 = DecisionTreeRegressionModel.load(model_path)
>>> model.numNodes == model2.numNodes
True
>>> model.depth == model2.depth
True
>>> model.transform(test1).head().variance
0.0

```

New in version 1.4.0.

setParams (*self*, *featuresCol*="features", *labelCol*="label", *predictionCol*="prediction", *maxDepth*=5, *maxBins*=32, *minInstancesPerNode*=1, *minInfoGain*=0.0, *maxMemoryInMB*=256, *cacheNodeIds*=False, *checkpointInterval*=10, *impurity*="variance", *seed*=None, *varianceCol*=None)
 Sets params for the DecisionTreeRegressor.

New in version 1.4.0.

class pyspark.ml.regression.**DecisionTreeRegressionModel** (*java_model*=None)
 Model fitted by *DecisionTreeRegressor*.

New in version 1.4.0.

featureImportances

Estimate of the importance of each feature.

This generalizes the idea of “Gini” importance to other losses, following the explanation of Gini importance from “Random Forests” documentation by Leo Breiman and Adele Cutler, and following the implementation from scikit-learn.

This feature importance is calculated as follows:

- $\text{importance}(\text{feature } j) = \text{sum (over nodes which split on feature } j) \text{ of the gain,}$
where gain is scaled by the number of instances passing through node
- Normalize importances for tree to sum to 1.

Note: Feature importance for single decision trees can have high variance due to correlated predictor variables. Consider using a *RandomForestRegressor* to determine feature importance instead.

New in version 2.0.0.

```
class pyspark.ml.regression.GBTRegressor (featuresCol='features',
                                           labelCol='label',      pre-
                                           dictionCol='prediction',
                                           maxDepth=5,  maxBins=32,
                                           minInstancesPerNode=1,
                                           minInfoGain=0.0,  maxMem-
                                           oryInMB=256,      cacheN-
                                           odeIds=False,  subsamplin-
                                           gRate=1.0,  checkpointInter-
                                           val=10,  lossType='squared',
                                           maxIter=20,  stepSize=0.1,
                                           seed=None,      impu-
                                           rity='variance',      feature-
                                           SubsetStrategy='all')
```

Gradient-Boosted Trees (GBTs) learning algorithm for regression. It supports both continuous and categorical features.

```
>>> from numpy import allclose
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> gbt = GBTRegressor(maxIter=5, maxDepth=2, seed=42)
>>> print(gbt.getImpurity())
variance
>>> print(gbt.getFeatureSubsetStrategy())
all
>>> model = gbt.fit(df)
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> model.numFeatures
1
>>> allclose(model.treeWeights, [1.0, 0.1, 0.1, 0.1, 0.1])
True
```

(continues on next page)

(continued from previous page)

```

>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
    ↪ "features"])
>>> model.transform(test0).head().prediction
0.0
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
    ↪)], ["features"])
>>> model.transform(test1).head().prediction
1.0
>>> gbtr_path = temp_path + "gbtr"
>>> gbt.save(gbtr_path)
>>> gbt2 = GBTRegressor.load(gbtr_path)
>>> gbt2.getMaxDepth()
2
>>> model_path = temp_path + "gbtr_model"
>>> model.save(model_path)
>>> model2 = GBTRegressionModel.load(model_path)
>>> model.featureImportances == model2.featureImportances
True
>>> model.treeWeights == model2.treeWeights
True
>>> model.trees
[DecisionTreeRegressionModel (uid=...) of depth..., ↪
    ↪ DecisionTreeRegressionModel...]
>>> validation = spark.createDataFrame([(0.0, Vectors.dense(-1.
    ↪ 0))],
    ... ["label", "features"])
>>> model.evaluateEachIteration(validation, "squared")
[0.0, 0.0, 0.0, 0.0, 0.0]

```

New in version 1.4.0.

getLossType()

Gets the value of lossType or its default value.

New in version 1.4.0.

setFeatureSubsetStrategy(value)

Sets the value of featureSubsetStrategy.

New in version 2.4.0.

setLossType(value)

Sets the value of lossType.

New in version 1.4.0.

```
setParams (self, featuresCol="features", labelCol="label", prediction-  
Col="prediction", maxDepth=5, maxBins=32, minInstances-  
PerNode=1, minInfoGain=0.0, maxMemoryInMB=256, cacheN-  
odeIds=False, subsamplingRate=1.0, checkpointInterval=10,  
lossType="squared", maxIter=20, stepSize=0.1, seed=None, im-  
purity="variance", featureSubsetStrategy="all")
```

Sets params for Gradient Boosted Tree Regression.

New in version 1.4.0.

```
class pyspark.ml.regression.GBTRegressionModel (java_model=None)  
Model fitted by GBTRegressor.
```

New in version 1.4.0.

```
evaluateEachIteration (dataset, loss)
```

Method to compute error or loss for every iteration of gradient boosting.

Parameters

- **dataset** – Test dataset to evaluate model on, where dataset is an instance of `pyspark.sql.DataFrame`
- **loss** – The loss function used to compute error. Supported options: squared, absolute

New in version 2.4.0.

featureImportances

Estimate of the importance of each feature.

Each feature's importance is the average of its importance across all trees in the ensemble. The importance vector is normalized to sum to 1. This method is suggested by Hastie et al. (Hastie, Tibshirani, Friedman. "The Elements of Statistical Learning, 2nd Edition." 2001.) and follows the implementation from scikit-learn.

See also:

DecisionTreeRegressionModel.featureImportances

New in version 2.0.0.

trees

These have null parent Estimators.

New in version 2.0.0.

Type Trees in this ensemble. Warning


```
class pyspark.ml.regression.GeneralizedLinearRegression (labelCol='label',  
                                                    fea-  
                                                    turesCol='features',  
                                                    predic-  
                                                    tion-  
                                                    Col='prediction',  
                                                    fam-  
                                                    ily='gaussian',  
                                                    link=None,  
                                                    fitInter-  
                                                    cept=True,  
                                                    max-  
                                                    Iter=25,  
                                                    tol=1e-  
                                                    06, reg-  
                                                    Param=0.0,  
                                                    weight-  
                                                    Col=None,  
                                                    solver='irls',  
                                                    linkPre-  
                                                    dic-  
                                                    tion-  
                                                    Col=None,  
                                                    vari-  
                                                    ance-  
                                                    Power=0.0,  
                                                    linkPower=None,  
                                                    offset-  
                                                    Col=None)
```

Note: Experimental

Generalized Linear Regression.

Fit a Generalized Linear Model specified by giving a symbolic description of the linear predictor (link function) and a description of the error distribution (family). It supports “gaussian”, “binomial”, “poisson”, “gamma” and “tweedie” as family. Valid link functions for each family is listed below. The first link function of each family is the default one.

- “gaussian” -> “identity”, “log”, “inverse”
- “binomial” -> “logit”, “probit”, “cloglog”
- “poisson” -> “log”, “identity”, “sqrt”
- “gamma” -> “inverse”, “identity”, “log”

- “tweedie” -> power link function specified through “linkPower”. The default link power in the tweedie family is 1 - variancePower.

See also:

GLM

```
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(0.0, 0.0)),
...     (1.0, Vectors.dense(1.0, 2.0)),
...     (2.0, Vectors.dense(0.0, 0.0)),
...     (2.0, Vectors.dense(1.0, 1.0)),], ["label", "features"])
>>> glr = GeneralizedLinearRegression(family="gaussian", link=
↳ "identity", linkPredictionCol="p")
>>> model = glr.fit(df)
>>> transformed = model.transform(df)
>>> abs(transformed.head().prediction - 1.5) < 0.001
True
>>> abs(transformed.head().p - 1.5) < 0.001
True
>>> model.coefficients
DenseVector([1.5..., -1.0...])
>>> model.numFeatures
2
>>> abs(model.intercept - 1.5) < 0.001
True
>>> glr_path = temp_path + "/glr"
>>> glr.save(glr_path)
>>> glr2 = GeneralizedLinearRegression.load(glr_path)
>>> glr.getFamily() == glr2.getFamily()
True
>>> model_path = temp_path + "/glr_model"
>>> model.save(model_path)
>>> model2 = GeneralizedLinearRegressionModel.load(model_path)
>>> model.intercept == model2.intercept
True
>>> model.coefficients[0] == model2.coefficients[0]
True
```

New in version 2.0.0.

getFamily()

Gets the value of family or its default value.

New in version 2.0.0.

getLink()

Gets the value of link or its default value.

New in version 2.0.0.

getLinkPower ()

Gets the value of linkPower or its default value.

New in version 2.2.0.

getLinkPredictionCol ()

Gets the value of linkPredictionCol or its default value.

New in version 2.0.0.

getOffsetCol ()

Gets the value of offsetCol or its default value.

New in version 2.3.0.

getVariancePower ()

Gets the value of variancePower or its default value.

New in version 2.2.0.

setFamily (value)

Sets the value of family.

New in version 2.0.0.

setLink (value)

Sets the value of link.

New in version 2.0.0.

setLinkPower (value)

Sets the value of linkPower.

New in version 2.2.0.

setLinkPredictionCol (value)

Sets the value of linkPredictionCol.

New in version 2.0.0.

setOffsetCol (value)

Sets the value of offsetCol.

New in version 2.3.0.

setParams (self, labelCol="label", featuresCol="features", predictionCol="prediction", family="gaussian", link=None, fitIntercept=True, maxIter=25, tol=1e-6, regParam=0.0, weightCol=None, solver="irls", linkPredictionCol=None, variancePower=0.0, linkPower=None, offsetCol=None)

Sets params for generalized linear regression.

New in version 2.0.0.

setVariancePower (*value*)

Sets the value of `variancePower`.

New in version 2.2.0.

class `pyspark.ml.regression.GeneralizedLinearRegressionModel` (*java_model=None*)

Note: Experimental

Model fitted by *GeneralizedLinearRegression*.

New in version 2.0.0.

coefficients

Model coefficients.

New in version 2.0.0.

evaluate (*dataset*)

Evaluates the model on a test dataset.

Parameters **dataset** – Test dataset to evaluate model on, where dataset is an instance of `pyspark.sql.DataFrame`

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.0.0.

intercept

Model intercept.

New in version 2.0.0.

summary

Gets summary (e.g. residuals, deviance, pValues) of model on training set. An exception is thrown if *trainingSummary* is *None*.

New in version 2.0.0.

class `pyspark.ml.regression.GeneralizedLinearRegressionSummary` (*java_obj=None*)

Note: Experimental

Generalized linear regression results evaluated on a dataset.

New in version 2.0.0.

aic

Akaike's "An Information Criterion"(AIC) for the fitted model.

New in version 2.0.0.

degreesOfFreedom

Degrees of freedom.

New in version 2.0.0.

deviance

The deviance for the fitted model.

New in version 2.0.0.

dispersion

The dispersion of the fitted model. It is taken as 1.0 for the "binomial" and "poisson" families, and otherwise estimated by the residual Pearson's Chi-Squared statistic (which is defined as sum of the squares of the Pearson residuals) divided by the residual degrees of freedom.

New in version 2.0.0.

nullDeviance

The deviance for the null model.

New in version 2.0.0.

numInstances

Number of instances in DataFrame predictions.

New in version 2.2.0.

predictionCol

Field in *predictions* which gives the predicted value of each instance. This is set to a new column name if the original model's *predictionCol* is not set.

New in version 2.0.0.

predictions

Predictions output by the model's *transform* method.

New in version 2.0.0.

rank

The numeric rank of the fitted linear model.

New in version 2.0.0.

residualDegreeOfFreedom

The residual degrees of freedom.

New in version 2.0.0.

residualDegreeOfFreedomNull

The residual degrees of freedom for the null model.

New in version 2.0.0.

residuals (*residualsType*='deviance')

Get the residuals of the fitted model by type.

Parameters **residualsType** – The type of residuals which should be returned. Supported options: deviance (default), pearson, working, and response.

New in version 2.0.0.

class pyspark.ml.regression.GeneralizedLinearRegressionTrainingSummary (*java_*

Note: Experimental

Generalized linear regression training results.

New in version 2.0.0.

coefficientStandardErrors

Standard error of estimated coefficients and intercept.

If `GeneralizedLinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

New in version 2.0.0.

numIterations

Number of training iterations.

New in version 2.0.0.

pValues

Two-sided p-value of estimated coefficients and intercept.

If `GeneralizedLinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

New in version 2.0.0.

solver

The numeric solver used for training.

New in version 2.0.0.

tValues

T-statistic of estimated coefficients and intercept.

If `GeneralizedLinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

New in version 2.0.0.

```
class pyspark.ml.regression.IsotonicRegression (featuresCol='features',  
                                              labelCol='label',  
                                              prediction-  
                                              Col='prediction',  
                                              weightCol=None,  
                                              isotonic=True,  
                                              featureIndex=0)
```

Currently implemented using parallelized pool adjacent violators algorithm. Only univariate (single feature) algorithm supported.

```
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> ir = IsotonicRegression()
>>> model = ir.fit(df)
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
    ↪ "features"])
>>> model.transform(test0).head().prediction
0.0
>>> model.boundaries
DenseVector([0.0, 1.0])
>>> ir_path = temp_path + "/ir"
>>> ir.save(ir_path)
>>> ir2 = IsotonicRegression.load(ir_path)
>>> ir2.getIsotonic()
True
>>> model_path = temp_path + "/ir_model"
>>> model.save(model_path)
>>> model2 = IsotonicRegressionModel.load(model_path)
>>> model.boundaries == model2.boundaries
True
>>> model.predictions == model2.predictions
True
```

New in version 1.6.0.

getFeatureIndex()

Gets the value of `featureIndex` or its default value.

getIsotonic()

Gets the value of `isotonic` or its default value.

setFeatureIndex(value)

Sets the value of `featureIndex`.

setIsotonic (*value*)

Sets the value of `isotonic`.

setParams (*featuresCol='features', labelCol='label', predictionCol='prediction',
weightCol=None, isotonic=True, featureIndex=0*)

`setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction",
weightCol=None, isotonic=True, featureIndex=0)`: Set the params for `IsotonicRegression`.

class `pyspark.ml.regression.IsotonicRegressionModel` (*java_model=None*)
Model fitted by *IsotonicRegression*.

New in version 1.6.0.

boundaries

Boundaries in increasing order for which predictions are known.

New in version 1.6.0.

predictions

Predictions associated with the boundaries at the same index, monotone because of isotonic regression.

New in version 1.6.0.

class `pyspark.ml.regression.LinearRegression` (*featuresCol='features',
labelCol='label',
predictionCol='prediction',
maxIter=100, regParam=0.0, elasticNetParam=0.0, tol=1e-06,
fitIntercept=True,
standardization=True,
solver='auto', weightCol=None, aggregationDepth=2,
loss='squaredError',
epsilon=1.35*)

Linear regression.

The learning objective is to minimize the specified loss function, with regularization. This supports two kinds of loss:

- `squaredError` (a.k.a squared loss)
- `huber` (a hybrid of squared error for relatively small errors and absolute error for relatively large ones, and we estimate the scale parameter from training data)

This supports multiple types of regularization:

- none (a.k.a. ordinary least squares)
- L2 (ridge regression)
- L1 (Lasso)
- L2 + L1 (elastic net)

Note: Fitting with huber loss only supports none and L2 regularization.

```
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, 2.0, Vectors.dense(1.0)),
...     (0.0, 2.0, Vectors.sparse(1, [], []))], ["label", "weight",
↪ "features"])
>>> lr = LinearRegression(maxIter=5, regParam=0.0, solver="normal",
↪ weightCol="weight")
>>> model = lr.fit(df)
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
↪ "features"])
>>> abs(model.transform(test0).head().prediction - (-1.0)) < 0.001
True
>>> abs(model.coefficients[0] - 1.0) < 0.001
True
>>> abs(model.intercept - 0.0) < 0.001
True
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
↪)], ["features"])
>>> abs(model.transform(test1).head().prediction - 1.0) < 0.001
True
>>> lr.setParams("vector")
Traceback (most recent call last):
...
TypeError: Method setParams forces keyword arguments.
>>> lr_path = temp_path + "/lr"
>>> lr.save(lr_path)
>>> lr2 = LinearRegression.load(lr_path)
>>> lr2.getMaxIter()
5
>>> model_path = temp_path + "/lr_model"
>>> model.save(model_path)
>>> model2 = LinearRegressionModel.load(model_path)
>>> model.coefficients[0] == model2.coefficients[0]
True
>>> model.intercept == model2.intercept
True
>>> model.numFeatures
```

(continues on next page)

(continued from previous page)

```
1
>>> model.write().format("pmml").save(model_path + "_2")
```

New in version 1.4.0.

getEpsilon()

Gets the value of epsilon or its default value.

New in version 2.3.0.

setEpsilon(value)

Sets the value of epsilon.

New in version 2.3.0.

setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", maxIter=100, regParam=0.0, elasticNetParam=0.0, tol=1e-6, fitIntercept=True, standardization=True, solver="auto", weightCol=None, aggregationDepth=2, loss="squaredError", epsilon=1.35)

Sets params for linear regression.

New in version 1.4.0.

class pyspark.ml.regression.**LinearRegressionModel**(java_model=None)
Model fitted by *LinearRegression*.

New in version 1.4.0.

coefficients

Model coefficients.

New in version 2.0.0.

evaluate(dataset)

Evaluates the model on a test dataset.

Parameters dataset – Test dataset to evaluate model on, where dataset is an instance of `pyspark.sql.DataFrame`

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.0.0.

intercept

Model intercept.

New in version 1.4.0.

scale

The value by which $\|y - X'w\|$ is scaled down when loss is “huber”, otherwise 1.0.

New in version 2.3.0.

summary

Gets summary (e.g. residuals, mse, r-squared) of model on training set. An exception is thrown if *trainingSummary* is *None*.

New in version 2.0.0.

class `pyspark.ml.regression.LinearRegressionSummary` (*java_obj=None*)

Note: Experimental

Linear regression results evaluated on a dataset.

New in version 2.0.0.

coefficientStandardErrors

Standard error of estimated coefficients and intercept. This value is only available when using the “normal” solver.

If `LinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

See also:

`LinearRegression.solver`

New in version 2.0.0.

degreesOfFreedom

Degrees of freedom.

New in version 2.2.0.

devianceResiduals

The weighted residuals, the usual residuals rescaled by the square root of the instance weights.

New in version 2.0.0.

explainedVariance

Returns the explained variance regression score. $\text{explainedVariance} = 1 - \frac{\text{variance}(y - \hat{y})}{\text{variance}(y)}$

See also:

[Wikipedia explain variation](#)

Note: This ignores instance weights (setting all to 1.0) from *LinearRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

featuresCol

Field in “predictions” which gives the features of each instance as a vector.

New in version 2.0.0.

labelCol

Field in “predictions” which gives the true label of each instance.

New in version 2.0.0.

meanAbsoluteError

Returns the mean absolute error, which is a risk function corresponding to the expected value of the absolute error loss or l1-norm loss.

Note: This ignores instance weights (setting all to 1.0) from *LinearRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

meanSquaredError

Returns the mean squared error, which is a risk function corresponding to the expected value of the squared error loss or quadratic loss.

Note: This ignores instance weights (setting all to 1.0) from *LinearRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

numInstances

Number of instances in DataFrame predictions

New in version 2.0.0.

pValues

Two-sided p-value of estimated coefficients and intercept. This value is only available when using the “normal” solver.

If `LinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

See also:

`LinearRegression.solver`

New in version 2.0.0.

predictionCol

Field in “predictions” which gives the predicted value of the label at each instance.

New in version 2.0.0.

predictions

Dataframe outputted by the model’s *transform* method.

New in version 2.0.0.

r2

Returns R^2 , the coefficient of determination.

See also:

[Wikipedia coefficient of determination](#)

Note: This ignores instance weights (setting all to 1.0) from *LinearRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

r2adj

Returns Adjusted R^2 , the adjusted coefficient of determination.

See also:

[Wikipedia coefficient of determination, Adjusted \$R^2\$](#)

Note: This ignores instance weights (setting all to 1.0) from *LinearRegression.weightCol*. This will change in later Spark versions.

New in version 2.4.0.

residuals

Residuals (label - predicted value)

New in version 2.0.0.

rootMeanSquaredError

Returns the root mean squared error, which is defined as the square root of the mean squared error.

Note: This ignores instance weights (setting all to 1.0) from

LinearRegression.weightCol. This will change in later Spark versions.

New in version 2.0.0.

tValues

T-statistic of estimated coefficients and intercept. This value is only available when using the “normal” solver.

If `LinearRegression.fitIntercept` is set to `True`, then the last element returned corresponds to the intercept.

See also:

`LinearRegression.solver`

New in version 2.0.0.

class `pyspark.ml.regression.LinearRegressionTrainingSummary` (*java_obj=None*)

Note: Experimental

Linear regression training results. Currently, the training summary ignores the training weights except for the objective trace.

New in version 2.0.0.

objectiveHistory

Objective function (scaled loss + regularization) at each iteration. This value is only available when using the “l-bfgs” solver.

See also:

`LinearRegression.solver`

New in version 2.0.0.

totalIterations

Number of training iterations until termination. This value is only available when using the “l-bfgs” solver.

See also:

`LinearRegression.solver`

New in version 2.0.0.

```
class pyspark.ml.regression.RandomForestRegressor (featuresCol='features',
                                                    label-
                                                    Col='label',
                                                    prediction-
                                                    Col='prediction',
                                                    maxDepth=5,
                                                    maxBins=32,
                                                    minInstances-
                                                    PerNode=1,
                                                    minInfo-
                                                    Gain=0.0,
                                                    maxMemory-
                                                    InMB=256,
                                                    cacheN-
                                                    odeIds=False,
                                                    checkpointInter-
                                                    val=10, impu-
                                                    rity='variance',
                                                    subsamplin-
                                                    gRate=1.0,
                                                    seed=None,
                                                    numTrees=20,
                                                    featureSub-
                                                    setStrat-
                                                    egy='auto')
```

Random Forest learning algorithm for regression. It supports both continuous and categorical features.

```
>>> from numpy import allclose
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> rf = RandomForestRegressor(numTrees=2, maxDepth=2, seed=42)
>>> model = rf.fit(df)
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> allclose(model.treeWeights, [1.0, 1.0])
True
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
    ↪ "features"])
>>> model.transform(test0).head().prediction
0.0
>>> model.numFeatures
1
```

(continues on next page)

(continued from previous page)

```

>>> model.trees
[DecisionTreeRegressionModel (uid=...) of depth...,
 ↳DecisionTreeRegressionModel...]
>>> model.getNumTrees
2
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
 ↳)], ["features"])
>>> model.transform(test1).head().prediction
0.5
>>> rfr_path = temp_path + "/rfr"
>>> rf.save(rfr_path)
>>> rf2 = RandomForestRegressor.load(rfr_path)
>>> rf2.getNumTrees()
2
>>> model_path = temp_path + "/rfr_model"
>>> model.save(model_path)
>>> model2 = RandomForestRegressionModel.load(model_path)
>>> model.featureImportances == model2.featureImportances
True

```

New in version 1.4.0.

setFeatureSubsetStrategy (*value*)

Sets the value of featureSubsetStrategy.

New in version 2.4.0.

setParams (*self*, *featuresCol*="features", *labelCol*="label", *predictionCol*="prediction", *maxDepth*=5, *maxBins*=32, *minInstancesPerNode*=1, *minInfoGain*=0.0, *maxMemoryInMB*=256, *cacheNodeIds*=False, *checkpointInterval*=10, *impurity*="variance", *subsamplingRate*=1.0, *seed*=None, *numTrees*=20, *featureSubsetStrategy*="auto")

Sets params for linear regression.

New in version 1.4.0.

class pyspark.ml.regression.**RandomForestRegressionModel** (*java_model*=None)
Model fitted by *RandomForestRegressor*.

New in version 1.4.0.

featureImportances

Estimate of the importance of each feature.

Each feature's importance is the average of its importance across all trees in the ensemble. The importance vector is normalized to sum to 1. This method is suggested by Hastie et al. (Hastie, Tibshirani, Friedman. "The Elements of Statistical Learning, 2nd Edition." 2001.) and follows the implementation from scikit-learn.

See also:

DecisionTreeRegressionModel.featureImportances

New in version 2.0.0.

trees

These have null parent Estimators.

New in version 2.0.0.

Type Trees in this ensemble. Warning

CLASSIFICATION API

```
class pyspark.ml.classification.LinearSVC(featuresCol='features',  
                                           labelCol='label', predictionCol='prediction', max-  
                                           Iter=100, regParam=0.0,  
                                           tol=1e-06, rawPredictionCol='rawPrediction',  
                                           fitIntercept=True, stan-  
                                           dardization=True, thresh-  
                                           old=0.0, weightCol=None,  
                                           aggregationDepth=2)
```

Note: Experimental

Linear SVM Classifier

This binary classifier optimizes the Hinge Loss using the OWLQN optimizer. Only supports L2 regularization currently.

```
>>> from pyspark.sql import Row
>>> from pyspark.ml.linalg import Vectors
>>> df = sc.parallelize([
...     Row(label=1.0, features=Vectors.dense(1.0, 1.0, 1.0)),
...     Row(label=0.0, features=Vectors.dense(1.0, 2.0, 3.0))]).
→toDF()
>>> svm = LinearSVC(maxIter=5, regParam=0.01)
>>> model = svm.fit(df)
>>> model.coefficients
DenseVector([0.0, -0.2792, -0.1833])
>>> model.intercept
1.0206118982229047
>>> model.numClasses
2
```

(continues on next page)

(continued from previous page)

```

>>> model.numFeatures
3
>>> test0 = sc.parallelize([Row(features=Vectors.dense(-1.0, -1.0, -1.0))]).toDF()
>>> result = model.transform(test0).head()
>>> result.prediction
1.0
>>> result.rawPrediction
DenseVector([-1.4831, 1.4831])
>>> svm_path = temp_path + "/svm"
>>> svm.save(svm_path)
>>> svm2 = LinearSVC.load(svm_path)
>>> svm2.getMaxIter()
5
>>> model_path = temp_path + "/svm_model"
>>> model.save(model_path)
>>> model2 = LinearSVCModel.load(model_path)
>>> model.coefficients[0] == model2.coefficients[0]
True
>>> model.intercept == model2.intercept
True

```

New in version 2.2.0.

setParams (*featuresCol='features', labelCol='label', predictionCol='prediction', maxIter=100, regParam=0.0, tol=1e-06, rawPredictionCol='rawPrediction', fitIntercept=True, standardization=True, threshold=0.0, weightCol=None, aggregationDepth=2*)
 setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", maxIter=100, regParam=0.0, tol=1e-6, rawPredictionCol="rawPrediction", fitIntercept=True, standardization=True, threshold=0.0, weightCol=None, aggregationDepth=2): Sets params for Linear SVM Classifier.

New in version 2.2.0.

class pyspark.ml.classification.**LinearSVCModel** (*java_model=None*)

Note: Experimental

Model fitted by LinearSVC.

New in version 2.2.0.

coefficients

Model coefficients of Linear SVM Classifier.

New in version 2.2.0.

intercept

Model intercept of Linear SVM Classifier.

New in version 2.2.0.

```
class pyspark.ml.classification.LogisticRegression (featuresCol='features',  
                                                    label-  
                                                    Col='label',  
                                                    prediction-  
                                                    Col='prediction',  
                                                    maxIter=100,  
                                                    regParam=0.0,  
                                                    elasticNet-  
                                                    Param=0.0,  
                                                    tol=1e-06,  
                                                    fitInter-  
                                                    cept=True,  
                                                    threshold=0.5,  
                                                    thresh-  
                                                    olds=None,  
                                                    probability-  
                                                    Col='probability',  
                                                    rawPrediction-  
                                                    Col='rawPrediction',  
                                                    standardiza-  
                                                    tion=True,  
                                                    weight-  
                                                    Col=None,  
                                                    aggregation-  
                                                    Depth=2,  
                                                    family='auto',  
                                                    lowerBound-  
                                                    sOnCoeffi-  
                                                    cients=None,  
                                                    upperBound-  
                                                    sOnCoeffi-  
                                                    cients=None,  
                                                    lowerBound-  
                                                    sOnInter-  
                                                    cepts=None,  
                                                    upperBound-  
                                                    sOnInter-  
                                                    cepts=None)
```

Logistic regression. This class supports multinomial logistic (softmax) and binomial logistic

regression.

```
>>> from pyspark.sql import Row
>>> from pyspark.ml.linalg import Vectors
>>> bdf = sc.parallelize([
...     Row(label=1.0, weight=1.0, features=Vectors.dense(0.0, 5.
... ↪0)),
...     Row(label=0.0, weight=2.0, features=Vectors.dense(1.0, 2.
... ↪0)),
...     Row(label=1.0, weight=3.0, features=Vectors.dense(2.0, 1.
... ↪0)),
...     Row(label=0.0, weight=4.0, features=Vectors.dense(3.0, 3.
... ↪0))]).toDF()
>>> blor = LogisticRegression(regParam=0.01, weightCol="weight")
>>> blorModel = blor.fit(bdf)
>>> blorModel.coefficients
DenseVector([-1.080..., -0.646...])
>>> blorModel.intercept
3.112...
>>> data_path = "data/mllib/sample_multiclass_classification_data.
... ↪txt"
>>> mdf = spark.read.format("libsvm").load(data_path)
>>> mlor = LogisticRegression(regParam=0.1, elasticNetParam=1.0,
... ↪family="multinomial")
>>> mlorModel = mlor.fit(mdf)
>>> mlorModel.coefficientMatrix
SparseMatrix(3, 4, [0, 1, 2, 3], [3, 2, 1], [1.87..., -2.75..., -0.
... ↪50...], 1)
>>> mlorModel.interceptVector
DenseVector([0.04..., -0.42..., 0.37...])
>>> test0 = sc.parallelize([Row(features=Vectors.dense(-1.0, 1.
... ↪0))]).toDF()
>>> result = blorModel.transform(test0).head()
>>> result.prediction
1.0
>>> result.probability
DenseVector([0.02..., 0.97...])
>>> result.rawPrediction
DenseVector([-3.54..., 3.54...])
>>> test1 = sc.parallelize([Row(features=Vectors.sparse(2, [0], [1.
... ↪0]))]).toDF()
>>> blorModel.transform(test1).head().prediction
1.0
>>> blor.setParams("vector")
Traceback (most recent call last):
...
TypeError: Method setParams forces keyword arguments.
```

(continues on next page)

(continued from previous page)

```

>>> lr_path = temp_path + "/lr"
>>> blor.save(lr_path)
>>> lr2 = LogisticRegression.load(lr_path)
>>> lr2.getRegParam()
0.01
>>> model_path = temp_path + "/lr_model"
>>> blorModel.save(model_path)
>>> model2 = LogisticRegressionModel.load(model_path)
>>> blorModel.coefficients[0] == model2.coefficients[0]
True
>>> blorModel.intercept == model2.intercept
True
>>> model2
LogisticRegressionModel: uid = ..., numClasses = 2, numFeatures = 2

```

New in version 1.3.0.

getFamily()

Gets the value of `family` or its default value.

New in version 2.1.0.

getLowerBoundsOnCoefficients()

Gets the value of `lowerBoundsOnCoefficients`

New in version 2.3.0.

getLowerBoundsOnIntercepts()

Gets the value of `lowerBoundsOnIntercepts`

New in version 2.3.0.

getThreshold()

Get threshold for binary classification.

If `thresholds` is set with length 2 (i.e., binary classification), this returns the equivalent threshold: $\frac{1}{1 + \frac{\text{thresholds}(0)}{\text{thresholds}(1)}}$. Otherwise, returns `threshold` if set or its default value if unset.

New in version 1.4.0.

getThresholds()

If `thresholds` is set, return its value. Otherwise, if `threshold` is set, return the equivalent thresholds for binary classification: (1-threshold, threshold). If neither are set, throw an error.

New in version 1.5.0.

getUpperBoundsOnCoefficients()

Gets the value of `upperBoundsOnCoefficients`

New in version 2.3.0.

getUpperBoundsOnIntercepts()

Gets the value of `upperBoundsOnIntercepts`

New in version 2.3.0.

setFamily(value)

Sets the value of `family`.

New in version 2.1.0.

setLowerBoundsOnCoefficients(value)

Sets the value of `lowerBoundsOnCoefficients`

New in version 2.3.0.

setLowerBoundsOnIntercepts(value)

Sets the value of `lowerBoundsOnIntercepts`

New in version 2.3.0.

setParams(*featuresCol='features', labelCol='label', predictionCol='prediction', maxIter=100, regParam=0.0, elasticNetParam=0.0, tol=1e-06, fitIntercept=True, threshold=0.5, thresholds=None, probabilityCol='probability', rawPredictionCol='rawPrediction', standardization=True, weightCol=None, aggregationDepth=2, family='auto', lowerBoundsOnCoefficients=None, upperBoundsOnCoefficients=None, lowerBoundsOnIntercepts=None, upperBoundsOnIntercepts=None*)

`setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", maxIter=100, regParam=0.0, elasticNetParam=0.0, tol=1e-6, fitIntercept=True, threshold=0.5, thresholds=None, probabilityCol="probability", rawPredictionCol="rawPrediction", standardization=True, weightCol=None, aggregationDepth=2, family="auto", lowerBoundsOnCoefficients=None, upperBoundsOnCoefficients=None, lowerBoundsOnIntercepts=None, upperBoundsOnIntercepts=None)`: Sets params for logistic regression. If the threshold and thresholds Params are both set, they must be equivalent.

New in version 1.3.0.

setThreshold(value)

Sets the value of `threshold`. Clears value of `thresholds` if it has been set.

New in version 1.4.0.

setThresholds(value)

Sets the value of `thresholds`. Clears value of `threshold` if it has been set.

New in version 1.5.0.

setUpperBoundsOnCoefficients (*value*)

Sets the value of `upperBoundsOnCoefficients`

New in version 2.3.0.

setUpperBoundsOnIntercepts (*value*)

Sets the value of `upperBoundsOnIntercepts`

New in version 2.3.0.

class `pyspark.ml.classification.LogisticRegressionModel` (*java_model=None*)

Model fitted by LogisticRegression.

New in version 1.3.0.

coefficientMatrix

Model coefficients.

New in version 2.1.0.

coefficients

Model coefficients of binomial logistic regression. An exception is thrown in the case of multinomial logistic regression.

New in version 2.0.0.

evaluate (*dataset*)

Evaluates the model on a test dataset.

Parameters **dataset** – Test dataset to evaluate model on, where dataset is an instance of `pyspark.sql.DataFrame`

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.0.0.

intercept

Model intercept of binomial logistic regression. An exception is thrown in the case of multinomial logistic regression.

New in version 1.4.0.

interceptVector

Model intercept.

New in version 2.1.0.

summary

Gets summary (e.g. accuracy/precision/recall, objective history, total iterations) of model trained on the training set. An exception is thrown if *trainingSummary* is *None*.

New in version 2.0.0.

class pyspark.ml.classification.LogisticRegressionSummary (*java_obj=None*)

Note: Experimental

Abstraction for Logistic Regression Results for a given model.

New in version 2.0.0.

accuracy

Returns accuracy. (equals to the total number of correctly classified instances out of the total number of instances.)

New in version 2.3.0.

fMeasureByLabel (*beta=1.0*)

Returns f-measure for each label (category).

New in version 2.3.0.

falsePositiveRateByLabel

Returns false positive rate for each label (category).

New in version 2.3.0.

featuresCol

Field in “predictions” which gives the features of each instance as a vector.

New in version 2.0.0.

labelCol

Field in “predictions” which gives the true label of each instance.

New in version 2.0.0.

labels

Returns the sequence of labels in ascending order. This order matches the order used in metrics which are specified as arrays over labels, e.g., truePositiveRateByLabel.

Note: In most cases, it will be values {0.0, 1.0, ..., numClasses-1}, However, if the training set is missing a label, then all of the arrays over labels (e.g., from truePositiveRateByLabel) will be of length numClasses-1 instead of the expected numClasses.

New in version 2.3.0.

precisionByLabel

Returns precision for each label (category).

New in version 2.3.0.

predictionCol

Field in “predictions” which gives the prediction of each class.

New in version 2.3.0.

predictions

Dataframe outputted by the model’s *transform* method.

New in version 2.0.0.

probabilityCol

Field in “predictions” which gives the probability of each class as a vector.

New in version 2.0.0.

recallByLabel

Returns recall for each label (category).

New in version 2.3.0.

truePositiveRateByLabel

Returns true positive rate for each label (category).

New in version 2.3.0.

weightedFMeasure (*beta=1.0*)

Returns weighted averaged f-measure.

New in version 2.3.0.

weightedFalsePositiveRate

Returns weighted false positive rate.

New in version 2.3.0.

weightedPrecision

Returns weighted averaged precision.

New in version 2.3.0.

weightedRecall

Returns weighted averaged recall. (equals to precision, recall and f-measure)

New in version 2.3.0.

weightedTruePositiveRate

Returns weighted true positive rate. (equals to precision, recall and f-measure)

New in version 2.3.0.

class pyspark.ml.classification.LogisticRegressionTrainingSummary (*java_obj=Non*

Note: Experimental

Abstraction for multinomial Logistic Regression Training results. Currently, the training summary ignores the training weights except for the objective trace.

New in version 2.0.0.

objectiveHistory

Objective function (scaled loss + regularization) at each iteration.

New in version 2.0.0.

totalIterations

Number of training iterations until termination.

New in version 2.0.0.

```
class pyspark.ml.classification.BinaryLogisticRegressionSummary (java_obj=None)
```

Note: Experimental

Binary Logistic regression results for a given model.

New in version 2.0.0.

areaUnderROC

Computes the area under the receiver operating characteristic (ROC) curve.

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

fMeasureByThreshold

Returns a dataframe with two fields (threshold, F-Measure) curve with beta = 1.0.

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

pr

Returns the precision-recall curve, which is a Dataframe containing two fields recall, precision with (0.0, 1.0) prepended to it.

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

precisionByThreshold

Returns a dataframe with two fields (threshold, precision) curve. Every possible probability obtained in transforming the dataset are used as thresholds used in calculating the precision.

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

recallByThreshold

Returns a dataframe with two fields (threshold, recall) curve. Every possible probability obtained in transforming the dataset are used as thresholds used in calculating the recall.

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

roc

Returns the receiver operating characteristic (ROC) curve, which is a Dataframe having two fields (FPR, TPR) with (0.0, 0.0) prepended and (1.0, 1.0) appended to it.

See also:

[Wikipedia reference](#)

Note: This ignores instance weights (setting all to 1.0) from *LogisticRegression.weightCol*. This will change in later Spark versions.

New in version 2.0.0.

class pyspark.ml.classification.**BinaryLogisticRegressionTrainingSummary** (*java_*

Note: Experimental

Binary Logistic regression training results for a given model.

New in version 2.0.0.

```
class pyspark.ml.classification.DecisionTreeClassifier (featuresCol='features',  
                                                    label-  
                                                    Col='label',  
                                                    predic-  
                                                    tion-  
                                                    Col='prediction',  
                                                    proba-  
                                                    bility-  
                                                    Col='probability',  
                                                    rawPre-  
                                                    diction-  
                                                    Col='rawPrediction',  
                                                    maxDepth=5,  
                                                    maxBins=32,  
                                                    minIn-  
                                                    stances-  
                                                    PerN-  
                                                    ode=1,  
                                                    minInfo-  
                                                    Gain=0.0,  
                                                    maxMem-  
                                                    ory-  
                                                    InMB=256,  
                                                    cacheN-  
                                                    odeIds=False,  
                                                    check-  
                                                    pointIn-  
                                                    ter-  
                                                    val=10,  
                                                    impu-  
                                                    rity='gini',  
                                                    seed=None)
```

Decision tree learning algorithm for classification. It supports both binary and multiclass labels, as well as both continuous and categorical features.

```
>>> from pyspark.ml.linalg import Vectors  
>>> from pyspark.ml.feature import StringIndexer  
>>> df = spark.createDataFrame([  
...     (1.0, Vectors.dense(1.0)),  
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])  
>>> stringIndexer = StringIndexer(inputCol="label", outputCol=  
    ↪ "indexed")
```

(continues on next page)

(continued from previous page)

```

>>> si_model = stringIndexer.fit(df)
>>> td = si_model.transform(df)
>>> dt = DecisionTreeClassifier(maxDepth=2, labelCol="indexed")
>>> model = dt.fit(td)
>>> model.numNodes
3
>>> model.depth
1
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> model.numFeatures
1
>>> model.numClasses
2
>>> print(model.toDebugString)
DecisionTreeClassificationModel (uid=...) of depth 1 with 3 nodes..
  ↳ .
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
  ↳ "features"])
>>> result = model.transform(test0).head()
>>> result.prediction
0.0
>>> result.probability
DenseVector([1.0, 0.0])
>>> result.rawPrediction
DenseVector([1.0, 0.0])
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
  ↳ )], ["features"])
>>> model.transform(test1).head().prediction
1.0

```

```

>>> dtc_path = temp_path + "/dtc"
>>> dt.save(dtc_path)
>>> dt2 = DecisionTreeClassifier.load(dtc_path)
>>> dt2.getMaxDepth()
2
>>> model_path = temp_path + "/dtc_model"
>>> model.save(model_path)
>>> model2 = DecisionTreeClassificationModel.load(model_path)
>>> model.featureImportances == model2.featureImportances
True

```

New in version 1.4.0.

```
setParams (self, featuresCol="features", labelCol="label", prediction-  
Col="prediction", probabilityCol="probability", rawPrediction-  
Col="rawPrediction", maxDepth=5, maxBins=32, minInstances-  
PerNode=1, minInfoGain=0.0, maxMemoryInMB=256, cacheN-  
odeIds=False, checkpointInterval=10, impurity="gini", seed=None)
```

Sets params for the DecisionTreeClassifier.

New in version 1.4.0.

```
class pyspark.ml.classification.DecisionTreeClassificationModel (java_model=None)  
Model fitted by DecisionTreeClassifier.
```

New in version 1.4.0.

featureImportances

Estimate of the importance of each feature.

This generalizes the idea of “Gini” importance to other losses, following the explanation of Gini importance from “Random Forests” documentation by Leo Breiman and Adele Cutler, and following the implementation from scikit-learn.

This feature importance is calculated as follows:

- $\text{importance}(\text{feature } j) = \text{sum (over nodes which split on feature } j) \text{ of the gain, where gain is scaled by the number of instances passing through node}$
- Normalize importances for tree to sum to 1.

Note: Feature importance for single decision trees can have high variance due to correlated predictor variables. Consider using a [*RandomForestClassifier*](#) to determine feature importance instead.

New in version 2.0.0.


```
class pyspark.ml.classification.GBTClassifier(featuresCol='features',
                                              labelCol='label',
                                              prediction-
                                              Col='prediction',
                                              maxDepth=5,
                                              maxBins=32, minIn-
                                              stancesPerNode=1,
                                              minInfoGain=0.0,
                                              maxMemory-
                                              InMB=256, cacheN-
                                              odeIds=False, check-
                                              pointInterval=10,
                                              lossType='logistic',
                                              maxIter=20, step-
                                              Size=0.1, seed=None,
                                              subsamplingRate=1.0,
                                              featureSubsetStrat-
                                              egy='all')
```

Gradient-Boosted Trees (GBTs) learning algorithm for classification. It supports binary labels, as well as both continuous and categorical features.

The implementation is based upon: J.H. Friedman. “Stochastic Gradient Boosting.” 1999.

Notes on Gradient Boosting vs. TreeBoost: - This implementation is for Stochastic Gradient Boosting, not for TreeBoost. - Both algorithms learn tree ensembles by minimizing loss functions. - TreeBoost (Friedman, 1999) additionally modifies the outputs at tree leaf nodes based on the loss function, whereas the original gradient boosting method does not. - We expect to implement TreeBoost in the future: [SPARK-4240](#)

Note: Multiclass labels are not currently supported.

```
>>> from numpy import allclose
>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.feature import StringIndexer
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> stringIndexer = StringIndexer(inputCol="label", outputCol=
    ↪ "indexed")
>>> si_model = stringIndexer.fit(df)
>>> td = si_model.transform(df)
>>> gbt = GBTCClassifier(maxIter=5, maxDepth=2, labelCol="indexed",
    ↪ seed=42)
>>> gbt.getFeatureSubsetStrategy()
'all'
```

(continues on next page)

(continued from previous page)

```

>>> model = gbt.fit(td)
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> allclose(model.treeWeights, [1.0, 0.1, 0.1, 0.1, 0.1])
True
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
    ↪ "features"])
>>> model.transform(test0).head().prediction
0.0
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
    ↪)], ["features"])
>>> model.transform(test1).head().prediction
1.0
>>> model.totalNumNodes
15
>>> print(model.toDebugString)
GBTCClassificationModel (uid=...)...with 5 trees...
>>> gbtc_path = temp_path + "gbtc"
>>> gbt.save(gbtc_path)
>>> gbt2 = GBTCClassifier.load(gbtc_path)
>>> gbt2.getMaxDepth()
2
>>> model_path = temp_path + "gbtc_model"
>>> model.save(model_path)
>>> model2 = GBTCClassificationModel.load(model_path)
>>> model.featureImportances == model2.featureImportances
True
>>> model.treeWeights == model2.treeWeights
True
>>> model.trees
[DecisionTreeRegressionModel (uid=...) of depth...,
    ↪ DecisionTreeRegressionModel...]
>>> validation = spark.createDataFrame([(0.0, Vectors.dense(-1.0),
    ↪)],
    ..., ["indexed", "features"])
>>> model.evaluateEachIteration(validation)
[0.25..., 0.23..., 0.21..., 0.19..., 0.18...]
>>> model.numClasses
2

```

New in version 1.4.0.

getLossType()

Gets the value of lossType or its default value.

New in version 1.4.0.

setFeatureSubsetStrategy (*value*)
Sets the value of `featureSubsetStrategy`.

New in version 2.4.0.

setLossType (*value*)
Sets the value of `lossType`.

New in version 1.4.0.

setParams (*self*, *featuresCol*="features", *labelCol*="label", *predictionCol*="prediction", *maxDepth*=5, *maxBins*=32, *minInstancesPerNode*=1, *minInfoGain*=0.0, *maxMemoryInMB*=256, *cacheNodeIds*=False, *checkpointInterval*=10, *lossType*="logistic", *maxIter*=20, *stepSize*=0.1, *seed*=None, *subsamplingRate*=1.0, *featureSubsetStrategy*="all")

Sets params for Gradient Boosted Tree Classification.

New in version 1.4.0.

class `pyspark.ml.classification.GBTClassificationModel` (*java_model*=None)
Model fitted by GBTClassifier.

New in version 1.4.0.

evaluateEachIteration (*dataset*)
Method to compute error or loss for every iteration of gradient boosting.

Parameters **dataset** – Test dataset to evaluate model on, where dataset is an instance of `pyspark.sql.DataFrame`

New in version 2.4.0.

featureImportances
Estimate of the importance of each feature.

Each feature's importance is the average of its importance across all trees in the ensemble. The importance vector is normalized to sum to 1. This method is suggested by Hastie et al. (Hastie, Tibshirani, Friedman. "The Elements of Statistical Learning, 2nd Edition." 2001.) and follows the implementation from scikit-learn.

See also:

[*DecisionTreeClassificationModel.featureImportances*](#)

New in version 2.0.0.

trees
These have null parent Estimators.

New in version 2.0.0.

Type Trees in this ensemble. Warning

```
class pyspark.ml.classification.RandomForestClassifier (featuresCol='features',  
                                                    label-  
                                                    Col='label',  
                                                    predic-  
                                                    tion-  
                                                    Col='prediction',  
                                                    proba-  
                                                    bility-  
                                                    Col='probability',  
                                                    rawPre-  
                                                    diction-  
                                                    Col='rawPrediction',  
                                                    maxDepth=5,  
                                                    maxBins=32,  
                                                    minIn-  
                                                    stances-  
                                                    PerN-  
                                                    ode=1,  
                                                    minInfo-  
                                                    Gain=0.0,  
                                                    maxMem-  
                                                    ory-  
                                                    InMB=256,  
                                                    cacheN-  
                                                    odeIds=False,  
                                                    check-  
                                                    pointIn-  
                                                    ter-  
                                                    val=10,  
                                                    impu-  
                                                    rity='gini',  
                                                    numTrees=20,  
                                                    feature-  
                                                    Subset-  
                                                    Strat-  
                                                    egy='auto',  
                                                    seed=None,  
                                                    subsam-  
                                                    plin-  
                                                    gRate=1.0)
```

Random Forest learning algorithm for classification. It supports both binary and multiclass labels, as well as both continuous and categorical features.

```
>>> import numpy
```

(continues on next page)

(continued from previous page)

```

>>> from numpy import allclose
>>> from pyspark.ml.linalg import Vectors
>>> from pyspark.ml.feature import StringIndexer
>>> df = spark.createDataFrame([
...     (1.0, Vectors.dense(1.0)),
...     (0.0, Vectors.sparse(1, [], []))], ["label", "features"])
>>> stringIndexer = StringIndexer(inputCol="label", outputCol=
↳ "indexed")
>>> si_model = stringIndexer.fit(df)
>>> td = si_model.transform(df)
>>> rf = RandomForestClassifier(numTrees=3, maxDepth=2, labelCol=
↳ "indexed", seed=42)
>>> model = rf.fit(td)
>>> model.featureImportances
SparseVector(1, {0: 1.0})
>>> allclose(model.treeWeights, [1.0, 1.0, 1.0])
True
>>> test0 = spark.createDataFrame([(Vectors.dense(-1.0),)], [
↳ "features"])
>>> result = model.transform(test0).head()
>>> result.prediction
0.0
>>> numpy.argmax(result.probability)
0
>>> numpy.argmax(result.rawPrediction)
0
>>> test1 = spark.createDataFrame([(Vectors.sparse(1, [0], [1.0]),
↳ )], ["features"])
>>> model.transform(test1).head().prediction
1.0
>>> model.trees
[DecisionTreeClassificationModel (uid=...) of depth...,
↳ DecisionTreeClassificationModel...]
>>> rfc_path = temp_path + "/rfc"
>>> rf.save(rfc_path)
>>> rf2 = RandomForestClassifier.load(rfc_path)
>>> rf2.getNumTrees()
3
>>> model_path = temp_path + "/rfc_model"
>>> model.save(model_path)
>>> model2 = RandomForestClassificationModel.load(model_path)
>>> model.featureImportances == model2.featureImportances
True

```

New in version 1.4.0.

setFeatureSubsetStrategy (*value*)

Sets the value of `featureSubsetStrategy`.

New in version 2.4.0.

setParams (*self*, *featuresCol*="features", *labelCol*="label", *predictionCol*="prediction", *probabilityCol*="probability", *rawPredictionCol*="rawPrediction", *maxDepth*=5, *maxBins*=32, *minInstancesPerNode*=1, *minInfoGain*=0.0, *maxMemoryInMB*=256, *cacheNodeIds*=False, *checkpointInterval*=10, *seed*=None, *impurity*="gini", *numTrees*=20, *featureSubsetStrategy*="auto", *subsamplingRate*=1.0)

Sets params for linear classification.

New in version 1.4.0.

class `pyspark.ml.classification.RandomForestClassificationModel` (*java_model*=None)
Model fitted by `RandomForestClassifier`.

New in version 1.4.0.

featureImportances

Estimate of the importance of each feature.

Each feature's importance is the average of its importance across all trees in the ensemble. The importance vector is normalized to sum to 1. This method is suggested by Hastie et al. (Hastie, Tibshirani, Friedman. "The Elements of Statistical Learning, 2nd Edition." 2001.) and follows the implementation from `scikit-learn`.

See also:

`DecisionTreeClassificationModel.featureImportances`

New in version 2.0.0.

trees

These have null parent Estimators.

New in version 2.0.0.

Type Trees in this ensemble. Warning

class `pyspark.ml.classification.NaiveBayes` (*featuresCol*='features', *labelCol*='label', *predictionCol*='prediction', *probabilityCol*='probability', *rawPredictionCol*='rawPrediction', *smoothing*=1.0, *modelType*='multinomial', *thresholds*=None, *weightCol*=None)

Naive Bayes Classifiers. It supports both Multinomial and Bernoulli NB. [Multinomial NB](#)

can handle finitely supported discrete data. For example, by converting documents into TF-IDF vectors, it can be used for document classification. By making every vector a binary (0/1) data, it can also be used as [Bernoulli NB](#). The input feature values must be nonnegative.

```
>>> from pyspark.sql import Row
>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     Row(label=0.0, weight=0.1, features=Vectors.dense([0.0, 0.
... ↪0])),
...     Row(label=0.0, weight=0.5, features=Vectors.dense([0.0, 1.
... ↪0])),
...     Row(label=1.0, weight=1.0, features=Vectors.dense([1.0, 0.
... ↪0]))])
>>> nb = NaiveBayes(smoothing=1.0, modelType="multinomial",
... ↪weightCol="weight")
>>> model = nb.fit(df)
>>> model.pi
DenseVector([-0.81..., -0.58...])
>>> model.theta
DenseMatrix(2, 2, [-0.91..., -0.51..., -0.40..., -1.09...], 1)
>>> test0 = sc.parallelize([Row(features=Vectors.dense([1.0, 0.
... ↪0]))]).toDF()
>>> result = model.transform(test0).head()
>>> result.prediction
1.0
>>> result.probability
DenseVector([0.32..., 0.67...])
>>> result.rawPrediction
DenseVector([-1.72..., -0.99...])
>>> test1 = sc.parallelize([Row(features=Vectors.sparse(2, [0], [1.
... ↪0]))]).toDF()
>>> model.transform(test1).head().prediction
1.0
>>> nb_path = temp_path + "/nb"
>>> nb.save(nb_path)
>>> nb2 = NaiveBayes.load(nb_path)
>>> nb2.getSmoothing()
1.0
>>> model_path = temp_path + "/nb_model"
>>> model.save(model_path)
>>> model2 = NaiveBayesModel.load(model_path)
>>> model.pi == model2.pi
True
>>> model.theta == model2.theta
True
>>> nb = nb.setThresholds([0.01, 10.00])
>>> model3 = nb.fit(df)
```

(continues on next page)

(continued from previous page)

```
>>> result = model3.transform(test0).head()
>>> result.prediction
0.0
```

New in version 1.5.0.

getModelType()

Gets the value of `modelType` or its default value.

New in version 1.5.0.

getSmoothing()

Gets the value of `smoothing` or its default value.

New in version 1.5.0.

setModelType(value)

Sets the value of `modelType`.

New in version 1.5.0.

setParams(*self*, *featuresCol*="features", *labelCol*="label", *predictionCol*="prediction", *probabilityCol*="probability", *rawPredictionCol*="rawPrediction", *smoothing*=1.0, *modelType*="multinomial", *thresholds*=None, *weightCol*=None)

Sets params for Naive Bayes.

New in version 1.5.0.

setSmoothing(value)

Sets the value of `smoothing`.

New in version 1.5.0.

class pyspark.ml.classification.NaiveBayesModel(*java_model*=None)

Model fitted by NaiveBayes.

New in version 1.5.0.

pi

log of class priors.

New in version 2.0.0.

theta

log of class conditional probabilities.

New in version 2.0.0.


```

class pyspark.ml.classification.MultilayerPerceptronClassifier (featuresCol='features',
    la-
    bel-
    Col='label',
    pre-
    dic-
    tion-
    Col='prediction',
    max-
    Iter=100,
    tol=1e-
    06,
    seed=None,
    lay-
    ers=None,
    block-
    Size=128,
    step-
    Size=0.03,
    solver='l-
    bfgs',
    ini-
    tial-
    Weights=None,
    prob-
    a-
    bil-
    i-
    ty-
    Col='probability',
    raw-
    Pre-
    dic-
    tion-
    Col='rawPrediction')

```

Classifier trainer based on the Multilayer Perceptron. Each layer has sigmoid activation function, output layer has softmax. Number of inputs has to be equal to the size of feature vectors. Number of outputs has to be equal to the total number of labels.

```

>>> from pyspark.ml.linalg import Vectors
>>> df = spark.createDataFrame([
...     (0.0, Vectors.dense([0.0, 0.0])),
...     (1.0, Vectors.dense([0.0, 1.0])),
...     (1.0, Vectors.dense([1.0, 0.0])),
...     (0.0, Vectors.dense([1.0, 1.0])), ["label", "features"]])

```

(continues on next page)

(continued from previous page)

```

>>> mlp = MultilayerPerceptronClassifier(maxIter=100, layers=[2, 2,
↪ 2], blockSize=1, seed=123)
>>> model = mlp.fit(df)
>>> model.layers
[2, 2, 2]
>>> model.weights.size
12
>>> testDF = spark.createDataFrame([
...     (Vectors.dense([1.0, 0.0]),),
...     (Vectors.dense([0.0, 0.0]),)], ["features"])
>>> model.transform(testDF).select("features", "prediction").show()
+-----+-----+
| features|prediction|
+-----+-----+
|[1.0,0.0]|      1.0|
|[0.0,0.0]|      0.0|
+-----+-----+
...
>>> mlp_path = temp_path + "/mlp"
>>> mlp.save(mlp_path)
>>> mlp2 = MultilayerPerceptronClassifier.load(mlp_path)
>>> mlp2.getBlockSize()
1
>>> model_path = temp_path + "/mlp_model"
>>> model.save(model_path)
>>> model2 = MultilayerPerceptronClassificationModel.load(model_
↪ path)
>>> model.layers == model2.layers
True
>>> model.weights == model2.weights
True
>>> mlp2 = mlp2.setInitialWeights(list(range(0, 12)))
>>> model3 = mlp2.fit(df)
>>> model3.weights != model2.weights
True
>>> model3.layers == model.layers
True

```

New in version 1.6.0.

getBlockSize()

Gets the value of blockSize or its default value.

New in version 1.6.0.

getInitialWeights()

Gets the value of initialWeights or its default value.

New in version 2.0.0.

getLayers()

Gets the value of layers or its default value.

New in version 1.6.0.

getStepSize()

Gets the value of stepSize or its default value.

New in version 2.0.0.

setBlockSize(value)

Sets the value of blockSize.

New in version 1.6.0.

setInitialWeights(value)

Sets the value of initialWeights.

New in version 2.0.0.

setLayers(value)

Sets the value of layers.

New in version 1.6.0.

setParams(*featuresCol='features', labelCol='label', predictionCol='prediction', maxIter=100, tol=1e-06, seed=None, layers=None, blockSize=128, stepSize=0.03, solver='l-bfgs', initialWeights=None, probabilityCol='probability', rawPredictionCol='rawPrediction'*)
setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", maxIter=100, tol=1e-6, seed=None, layers=None, blockSize=128, stepSize=0.03, solver="l-bfgs", initialWeights=None, probabilityCol="probability", rawPredictionCol="rawPrediction"): Sets params for MultilayerPerceptronClassifier.

New in version 1.6.0.

setStepSize(value)

Sets the value of stepSize.

New in version 2.0.0.

class pyspark.ml.classification.**MultilayerPerceptronClassificationModel** (*java_*
Model fitted by MultilayerPerceptronClassifier.

New in version 1.6.0.

layers

array of layer sizes including input and output layers.

New in version 1.6.0.

weights

the weights of layers.

New in version 2.0.0.

```
class pyspark.ml.classification.OneVsRest (featuresCol='features',  
                                           labelCol='label',           pre-  
                                           dictionCol='prediction',  
                                           classifier=None,       weight-  
                                           Col=None, parallelism=1)
```

Note: Experimental

Reduction of Multiclass Classification to Binary Classification. Performs reduction using one against all strategy. For a multiclass classification with k classes, train k models (one per class). Each example is scored against all k models and the model with highest score is picked to label the example.

```
>>> from pyspark.sql import Row
>>> from pyspark.ml.linalg import Vectors
>>> data_path = "data/mllib/sample_multiclass_classification_data.
↳txt"
>>> df = spark.read.format("libsvm").load(data_path)
>>> lr = LogisticRegression(regParam=0.01)
>>> ovr = OneVsRest(classifier=lr)
>>> model = ovr.fit(df)
>>> model.models[0].coefficients
DenseVector([0.5..., -1.0..., 3.4..., 4.2...])
>>> model.models[1].coefficients
DenseVector([-2.1..., 3.1..., -2.6..., -2.3...])
>>> model.models[2].coefficients
DenseVector([0.3..., -3.4..., 1.0..., -1.1...])
>>> [x.intercept for x in model.models]
[-2.7..., -2.5..., -1.3...]
>>> test0 = sc.parallelize([Row(features=Vectors.dense(-1.0, 0.0, 1.0, 1.0))]).toDF()
>>> model.transform(test0).head().prediction
0.0
>>> test1 = sc.parallelize([Row(features=Vectors.sparse(4, [0], [1.0]))]).toDF()
>>> model.transform(test1).head().prediction
2.0
>>> test2 = sc.parallelize([Row(features=Vectors.dense(0.5, 0.4, 0.3, 0.2))]).toDF()
>>> model.transform(test2).head().prediction
```

(continues on next page)

(continued from previous page)

```

0.0
>>> model_path = temp_path + "/ovr_model"
>>> model.save(model_path)
>>> model2 = OneVsRestModel.load(model_path)
>>> model2.transform(test0).head().prediction
0.0

```

New in version 2.0.0.

copy (*extra=None*)

Creates a copy of this instance with a randomly generated uid and some extra params. This creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over.

Parameters **extra** – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 2.0.0.

setParams (*featuresCol='features', labelCol='label', predictionCol='prediction', classifier=None, weightCol=None, parallelism=1*)
 setParams(self, featuresCol="features", labelCol="label", predictionCol="prediction", classifier=None, weightCol=None, parallelism=1): Sets params for OneVsRest.

New in version 2.0.0.

class pyspark.ml.classification.**OneVsRestModel** (*models*)

Note: Experimental

Model fitted by OneVsRest. This stores the models resulting from training k binary classifiers: one for each class. Each example is scored against all k models, and the model with the highest score is picked to label the example.

New in version 2.0.0.

copy (*extra=None*)

Creates a copy of this instance with a randomly generated uid and some extra params. This creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over.

Parameters **extra** – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 2.0.0.

CLUSTERING API

```
class pyspark.ml.clustering.BisectingKMeans (featuresCol='features',  
                                             prediction-  
                                             Col='prediction', max-  
                                             Iter=20, seed=None,  
                                             k=4, minDivisibleClus-  
                                             terSize=1.0, distance-  
                                             Measure='euclidean')
```

A bisecting k-means algorithm based on the paper “A comparison of document clustering techniques” by Steinbach, Karypis, and Kumar, with modification to fit Spark. The algorithm starts from a single cluster that contains all points. Iteratively it finds divisible clusters on the bottom level and bisects each of them using k-means, until there are k leaf clusters in total or no leaf clusters are divisible. The bisecting steps of clusters on the same level are grouped together to increase parallelism. If bisecting all divisible clusters on the bottom level would result more than k leaf clusters, larger clusters get higher priority.

```
>>> from pyspark.ml.linalg import Vectors  
>>> data = [(Vectors.dense([0.0, 0.0])), (Vectors.dense([1.0, 1.  
→0])),  
...         (Vectors.dense([9.0, 8.0])), (Vectors.dense([8.0, 9.  
→0]))]  
>>> df = spark.createDataFrame(data, ["features"])  
>>> bkm = BisectingKMeans(k=2, minDivisibleClusterSize=1.0)  
>>> model = bkm.fit(df)  
>>> centers = model.clusterCenters()  
>>> len(centers)  
2  
>>> model.computeCost(df)  
2.000...  
>>> model.hasSummary  
True  
>>> summary = model.summary  
>>> summary.k  
2
```

(continues on next page)

(continued from previous page)

```

>>> summary.clusterSizes
[2, 2]
>>> transformed = model.transform(df).select("features",
↳ "prediction")
>>> rows = transformed.collect()
>>> rows[0].prediction == rows[1].prediction
True
>>> rows[2].prediction == rows[3].prediction
True
>>> bkm_path = temp_path + "/bkm"
>>> bkm.save(bkm_path)
>>> bkm2 = BisectingKMeans.load(bkm_path)
>>> bkm2.getK()
2
>>> bkm2.getDistanceMeasure()
'euclidean'
>>> model_path = temp_path + "/bkm_model"
>>> model.save(model_path)
>>> model2 = BisectingKMeansModel.load(model_path)
>>> model2.hasSummary
False
>>> model.clusterCenters()[0] == model2.clusterCenters()[0]
array([ True,  True], dtype=bool)
>>> model.clusterCenters()[1] == model2.clusterCenters()[1]
array([ True,  True], dtype=bool)

```

New in version 2.0.0.

getDistanceMeasure()

Gets the value of *distanceMeasure* or its default value.

New in version 2.4.0.

getK()

Gets the value of *k* or its default value.

New in version 2.0.0.

getMinDivisibleClusterSize()

Gets the value of *minDivisibleClusterSize* or its default value.

New in version 2.0.0.

setDistanceMeasure(value)

Sets the value of *distanceMeasure*.

New in version 2.4.0.

setK(value)

Sets the value of `k`.

New in version 2.0.0.

setMinDivisibleClusterSize (*value*)

Sets the value of `minDivisibleClusterSize`.

New in version 2.0.0.

setParams (*self*, *featuresCol*="features", *predictionCol*="prediction", *maxIter*=20, *seed*=None, *k*=4, *minDivisibleClusterSize*=1.0, *distanceMeasure*="euclidean")

Sets params for `BisectingKMeans`.

New in version 2.0.0.

class `pyspark.ml.clustering.BisectingKMeansModel` (*java_model*=None)

Model fitted by `BisectingKMeans`.

New in version 2.0.0.

clusterCenters ()

Get the cluster centers, represented as a list of NumPy arrays.

New in version 2.0.0.

computeCost (*dataset*)

Computes the sum of squared distances between the input points and their corresponding cluster centers.

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.1.0.

summary

Gets summary (e.g. cluster assignments, cluster sizes) of the model trained on the training set. An exception is thrown if no summary exists.

New in version 2.1.0.

class `pyspark.ml.clustering.BisectingKMeansSummary` (*java_obj*=None)

Note: Experimental

Bisecting KMeans clustering results for a given model.

New in version 2.1.0.

```
class pyspark.ml.clustering.KMeans (featuresCol='features',    prediction-
                                   Col='prediction', k=2, initMode='k-
                                   means||', initSteps=2, tol=0.0001,
                                   maxIter=20, seed=None, distance-
                                   Measure='euclidean')
```

K-means clustering with a k-means++ like initialization mode (the k-means|| algorithm by Bahmani et al).

```
>>> from pyspark.ml.linalg import Vectors
>>> data = [(Vectors.dense([0.0, 0.0]),), (Vectors.dense([1.0, 1.
↪0]),),
...         (Vectors.dense([9.0, 8.0]),), (Vectors.dense([8.0, 9.
↪0]),)]
>>> df = spark.createDataFrame(data, ["features"])
>>> kmeans = KMeans(k=2, seed=1)
>>> model = kmeans.fit(df)
>>> centers = model.clusterCenters()
>>> len(centers)
2
>>> model.computeCost(df)
2.000...
>>> transformed = model.transform(df).select("features",
↪"prediction")
>>> rows = transformed.collect()
>>> rows[0].prediction == rows[1].prediction
True
>>> rows[2].prediction == rows[3].prediction
True
>>> model.hasSummary
True
>>> summary = model.summary
>>> summary.k
2
>>> summary.clusterSizes
[2, 2]
>>> summary.trainingCost
2.000...
>>> kmeans_path = temp_path + "/kmeans"
>>> kmeans.save(kmeans_path)
>>> kmeans2 = KMeans.load(kmeans_path)
>>> kmeans2.getK()
2
>>> model_path = temp_path + "/kmeans_model"
>>> model.save(model_path)
>>> model2 = KMeansModel.load(model_path)
>>> model2.hasSummary
False
```

(continues on next page)

(continued from previous page)

```
>>> model.clusterCenters()[0] == model2.clusterCenters()[0]
array([ True,  True], dtype=bool)
>>> model.clusterCenters()[1] == model2.clusterCenters()[1]
array([ True,  True], dtype=bool)
```

New in version 1.5.0.

getDistanceMeasure()

Gets the value of *distanceMeasure*

New in version 2.4.0.

getInitMode()

Gets the value of *initMode*

New in version 1.5.0.

getInitSteps()

Gets the value of *initSteps*

New in version 1.5.0.

getK()

Gets the value of *k*

New in version 1.5.0.

setDistanceMeasure(value)

Sets the value of *distanceMeasure*.

New in version 2.4.0.

setInitMode(value)

Sets the value of *initMode*.

New in version 1.5.0.

setInitSteps(value)

Sets the value of *initSteps*.

New in version 1.5.0.

setK(value)

Sets the value of *k*.

New in version 1.5.0.

setParams(self, featuresCol="features", predictionCol="prediction", k=2, initMode="k-means||", initSteps=2, tol=1e-4, maxIter=20, seed=None, distanceMeasure="euclidean")

Sets params for KMeans.

New in version 1.5.0.

class pyspark.ml.clustering.**KMeansModel** (*java_model=None*)

Model fitted by KMeans.

New in version 1.5.0.

clusterCenters ()

Get the cluster centers, represented as a list of NumPy arrays.

New in version 1.5.0.

computeCost (*dataset*)

Return the K-means cost (sum of squared distances of points to their nearest center) for this model on the given data.

..note:: Deprecated in 2.4.0. It will be removed in 3.0.0. Use ClusteringEvaluator instead.

You can also get the cost on the training dataset in the summary.

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.1.0.

summary

Gets summary (e.g. cluster assignments, cluster sizes) of the model trained on the training set. An exception is thrown if no summary exists.

New in version 2.1.0.

class pyspark.ml.clustering.**GaussianMixture** (*featuresCol='features',*
prediction-
Col='prediction',
k=2, probability-
Col='probability',
tol=0.01, maxIter=100,
seed=None)

GaussianMixture clustering. This class performs expectation maximization for multivariate Gaussian Mixture Models (GMMs). A GMM represents a composite distribution of independent Gaussian distributions with associated “mixing” weights specifying each’s contribution to the composite.

Given a set of sample points, this class will maximize the log-likelihood for a mixture of *k* Gaussians, iterating until the log-likelihood changes by less than *convergenceTol*, or until it has reached the max number of iterations. While this process is generally guaranteed to converge, it is not guaranteed to find a global optimum.

Note: For high-dimensional data (with many features), this algorithm may perform poorly. This is due to high-dimensional data (a) making it difficult to cluster at all (based on statisti-

cal/theoretical arguments) and (b) numerical issues with Gaussian distributions.

```
>>> from pyspark.ml.linalg import Vectors

>>> data = [(Vectors.dense([-0.1, -0.05 ]),),
...         (Vectors.dense([-0.01, -0.1]),),
...         (Vectors.dense([0.9, 0.8]),),
...         (Vectors.dense([0.75, 0.935]),),
...         (Vectors.dense([-0.83, -0.68]),),
...         (Vectors.dense([-0.91, -0.76]),)]
>>> df = spark.createDataFrame(data, ["features"])
>>> gm = GaussianMixture(k=3, tol=0.0001,
...                       maxIter=10, seed=10)
>>> model = gm.fit(df)
>>> model.hasSummary
True
>>> summary = model.summary
>>> summary.k
3
>>> summary.clusterSizes
[2, 2, 2]
>>> summary.logLikelihood
8.14636...
>>> weights = model.weights
>>> len(weights)
3
>>> model.gaussiansDF.select("mean").head()
Row(mean=DenseVector([0.825, 0.8675]))
>>> model.gaussiansDF.select("cov").head()
Row(cov=DenseMatrix(2, 2, [0.0056, -0.0051, -0.0051, 0.0046]),
    ↪False))
>>> transformed = model.transform(df).select("features",
    ↪"prediction")
>>> rows = transformed.collect()
>>> rows[4].prediction == rows[5].prediction
True
>>> rows[2].prediction == rows[3].prediction
True
>>> gmm_path = temp_path + "/gmm"
>>> gm.save(gmm_path)
>>> gm2 = GaussianMixture.load(gmm_path)
>>> gm2.getK()
3
>>> model_path = temp_path + "/gmm_model"
>>> model.save(model_path)
```

(continues on next page)

(continued from previous page)

```

>>> model2 = GaussianMixtureModel.load(model_path)
>>> model2.hasSummary
False
>>> model2.weights == model.weights
True
>>> model2.gaussiansDF.select("mean").head()
Row(mean=DenseVector([0.825, 0.8675]))
>>> model2.gaussiansDF.select("cov").head()
Row(cov=DenseMatrix(2, 2, [0.0056, -0.0051, -0.0051, 0.0046]),
↪False)

```

New in version 2.0.0.

getK()

Gets the value of k

New in version 2.0.0.

setK(value)

Sets the value of k .

New in version 2.0.0.

setParams (*self*, *featuresCol*="features", *predictionCol*="prediction", *k*=2, *probabilityCol*="probability", *tol*=0.01, *maxIter*=100, *seed*=None)

Sets params for GaussianMixture.

New in version 2.0.0.

class pyspark.ml.clustering.**GaussianMixtureModel** (*java_model*=None)

Model fitted by GaussianMixture.

New in version 2.0.0.

gaussiansDF

Retrieve Gaussian distributions as a DataFrame. Each row represents a Gaussian Distribution. The DataFrame has two columns: mean (Vector) and cov (Matrix).

New in version 2.0.0.

hasSummary

Indicates whether a training summary exists for this model instance.

New in version 2.1.0.

summary

Gets summary (e.g. cluster assignments, cluster sizes) of the model trained on the training set. An exception is thrown if no summary exists.

New in version 2.1.0.

weights

Weight for each Gaussian distribution in the mixture. This is a multinomial probability distribution over the k Gaussians, where `weights[i]` is the weight for Gaussian i , and weights sum to 1.

New in version 2.0.0.

```
class pyspark.ml.clustering.GaussianMixtureSummary (java_obj=None)
```

Note: Experimental

Gaussian mixture clustering results for a given model.

New in version 2.1.0.

logLikelihood

Total log-likelihood for this model on the given data.

New in version 2.2.0.

probability

DataFrame of probabilities of each cluster for each training data point.

New in version 2.1.0.

probabilityCol

Name for column of predicted probability of each cluster in *predictions*.

New in version 2.1.0.

```
class pyspark.ml.clustering.LDA (featuresCol='features',      maxIter=20,  
                                seed=None,      checkpointInterval=10,  
                                k=10,  optimizer='online',  learningOff-  
                                set=1024.0,  learningDecay=0.51,  sub-  
                                samplingRate=0.05,  optimizeDocConcen-  
                                tration=True,  docConcentration=None,  
                                topicConcentration=None,  topicDistri-  
                                butionCol='topicDistribution',  keepLas-  
                                tCheckpoint=True)
```

Latent Dirichlet Allocation (LDA), a topic model designed for text documents.

Terminology:

- “term” = “word”: an element of the vocabulary
- “token”: instance of a term appearing in a document
- “topic”: multinomial distribution over terms representing some concept
- “document”: one piece of text, corresponding to one row in the input data

Original LDA paper (journal version): Blei, Ng, and Jordan. “Latent Dirichlet Allocation.” JMLR, 2003.

Input data (featuresCol): LDA is given a collection of documents as input data, via the featuresCol parameter. Each document is specified as a `Vector` of length vocabSize, where each entry is the count for the corresponding term (word) in the document. Feature transformers such as `pyspark.ml.feature.Tokenizer` and `pyspark.ml.feature.CountVectorizer` can be useful for converting text to word count vectors.

```
>>> from pyspark.ml.linalg import Vectors, SparseVector
>>> from pyspark.ml.clustering import LDA
>>> df = spark.createDataFrame([[1, Vectors.dense([0.0, 1.0])],
...                             [2, SparseVector(2, {0: 1.0})]], ["id", "features"])
>>> lda = LDA(k=2, seed=1, optimizer="em")
>>> model = lda.fit(df)
>>> model.isDistributed()
True
>>> localModel = model.toLocal()
>>> localModel.isDistributed()
False
>>> model.vocabSize()
2
>>> model.describeTopics().show()
+-----+-----+-----+
|topic|termIndices|          termWeights|
+-----+-----+-----+
|    0|    [1, 0]| [0.50401530077160...|
|    1|    [0, 1]| [0.50401530077160...|
+-----+-----+-----+
...
>>> model.topicsMatrix()
DenseMatrix(2, 2, [0.496, 0.504, 0.504, 0.496], 0)
>>> lda_path = temp_path + "/lda"
>>> lda.save(lda_path)
>>> sameLDA = LDA.load(lda_path)
>>> distributed_model_path = temp_path + "/lda_distributed_model"
>>> model.save(distributed_model_path)
>>> sameModel = DistributedLDAModel.load(distributed_model_path)
>>> local_model_path = temp_path + "/lda_local_model"
>>> localModel.save(local_model_path)
>>> sameLocalModel = LocalLDAModel.load(local_model_path)
```

New in version 2.0.0.

getDocConcentration()

Gets the value of docConcentration or its default value.

New in version 2.0.0.

getK()

Gets the value of `k` or its default value.

New in version 2.0.0.

getKeepLastCheckpoint()

Gets the value of `keepLastCheckpoint` or its default value.

New in version 2.0.0.

getLearningDecay()

Gets the value of `learningDecay` or its default value.

New in version 2.0.0.

getLearningOffset()

Gets the value of `learningOffset` or its default value.

New in version 2.0.0.

getOptimizeDocConcentration()

Gets the value of `optimizeDocConcentration` or its default value.

New in version 2.0.0.

getOptimizer()

Gets the value of `optimizer` or its default value.

New in version 2.0.0.

getSubsamplingRate()

Gets the value of `subsamplingRate` or its default value.

New in version 2.0.0.

getTopicConcentration()

Gets the value of `topicConcentration` or its default value.

New in version 2.0.0.

getTopicDistributionCol()

Gets the value of `topicDistributionCol` or its default value.

New in version 2.0.0.

setDocConcentration(value)

Sets the value of `docConcentration`.

```
>>> algo = LDA().setDocConcentration([0.1, 0.2])
>>> algo.getDocConcentration()
[0.1..., 0.2...]
```

New in version 2.0.0.

setK(*value*)

Sets the value of k.

```
>>> algo = LDA().setK(10)
>>> algo.getK()
10
```

New in version 2.0.0.

setKeepLastCheckpoint(*value*)

Sets the value of keepLastCheckpoint.

```
>>> algo = LDA().setKeepLastCheckpoint(False)
>>> algo.getKeepLastCheckpoint()
False
```

New in version 2.0.0.

setLearningDecay(*value*)

Sets the value of learningDecay.

```
>>> algo = LDA().setLearningDecay(0.1)
>>> algo.getLearningDecay()
0.1...
```

New in version 2.0.0.

setLearningOffset(*value*)

Sets the value of learningOffset.

```
>>> algo = LDA().setLearningOffset(100)
>>> algo.getLearningOffset()
100.0
```

New in version 2.0.0.

setOptimizeDocConcentration(*value*)

Sets the value of optimizeDocConcentration.

```
>>> algo = LDA().setOptimizeDocConcentration(True)
>>> algo.getOptimizeDocConcentration()
True
```

New in version 2.0.0.

setOptimizer(*value*)

Sets the value of optimizer. Currently only support 'em' and 'online'.

```
>>> algo = LDA().setOptimizer("em")
>>> algo.getOptimizer()
'em'
```

New in version 2.0.0.

setParams (*self*, *featuresCol*="features", *maxIter*=20, *seed*=None, *checkpointInterval*=10, *k*=10, *optimizer*="online", *learningOffset*=1024.0, *learningDecay*=0.51, *subsamplingRate*=0.05, *optimizeDocConcentration*=True, *docConcentration*=None, *topicConcentration*=None, *topicDistributionCol*="topicDistribution", *keepLastCheckpoint*=True)

Sets params for LDA.

New in version 2.0.0.

setSubsamplingRate (*value*)

Sets the value of *subsamplingRate*.

```
>>> algo = LDA().setSubsamplingRate(0.1)
>>> algo.getSubsamplingRate()
0.1...
```

New in version 2.0.0.

setTopicConcentration (*value*)

Sets the value of *topicConcentration*.

```
>>> algo = LDA().setTopicConcentration(0.5)
>>> algo.getTopicConcentration()
0.5...
```

New in version 2.0.0.

setTopicDistributionCol (*value*)

Sets the value of *topicDistributionCol*.

```
>>> algo = LDA().setTopicDistributionCol("topicDistributionCol
↪")
>>> algo.getTopicDistributionCol()
'topicDistributionCol'
```

New in version 2.0.0.

class pyspark.ml.clustering.LDAModel (*java_model*=None)

Latent Dirichlet Allocation (LDA) model. This abstraction permits for different underlying representations, including local and distributed data structures.

New in version 2.0.0.

describeTopics (*maxTermsPerTopic=10*)

Return the topics described by their top-weighted terms.

New in version 2.0.0.

estimatedDocConcentration ()

Value for `LDA.docConcentration` estimated from data. If Online LDA was used and `LDA.optimizeDocConcentration` was set to `false`, then this returns the fixed (given) value for the `LDA.docConcentration` parameter.

New in version 2.0.0.

isDistributed ()

Indicates whether this instance is of type `DistributedLDAModel`

New in version 2.0.0.

logLikelihood (*dataset*)

Calculates a lower bound on the log likelihood of the entire corpus. See Equation (16) in the Online LDA paper (Hoffman et al., 2010).

WARNING: If this model is an instance of `DistributedLDAModel` (produced when `optimizer` is set to “em”), this involves collecting a large `topicsMatrix()` to the driver. This implementation may be changed in the future.

New in version 2.0.0.

logPerplexity (*dataset*)

Calculate an upper bound on perplexity. (Lower is better.) See Equation (16) in the Online LDA paper (Hoffman et al., 2010).

WARNING: If this model is an instance of `DistributedLDAModel` (produced when `optimizer` is set to “em”), this involves collecting a large `topicsMatrix()` to the driver. This implementation may be changed in the future.

New in version 2.0.0.

topicsMatrix ()

Inferred topics, where each topic is represented by a distribution over terms. This is a matrix of size `vocabSize` x `k`, where each column is a topic. No guarantees are given about the ordering of the topics.

WARNING: If this model is actually a `DistributedLDAModel` instance produced by the Expectation-Maximization (“em”) *optimizer*, then this method could involve collecting a large amount of data to the driver (on the order of `vocabSize` x `k`).

New in version 2.0.0.

vocabSize ()

Vocabulary size (number of terms or words in the vocabulary)

New in version 2.0.0.

class pyspark.ml.clustering.**LocalLDAModel** (*java_model=None*)

Local (non-distributed) model fitted by *LDA*. This model stores the inferred topics only; it does not store info about the training dataset.

New in version 2.0.0.

class pyspark.ml.clustering.**DistributedLDAModel** (*java_model=None*)

Distributed model fitted by *LDA*. This type of model is currently only produced by Expectation-Maximization (EM).

This model stores the inferred topics, the full training dataset, and the topic distribution for each training document.

New in version 2.0.0.

getCheckpointFiles ()

If using checkpointing and *LDA.keepLastCheckpoint* is set to true, then there may be saved checkpoint files. This method is provided so that users can manage those files.

Note: Removing the checkpoints can cause failures if a partition is lost and is needed by certain *DistributedLDAModel* methods. Reference counting will clean up the checkpoints when this model and derivative data go out of scope.

:return List of checkpoint files from training

New in version 2.0.0.

logPrior ()

Log probability of the current parameter estimate: $\log P(\text{topics, topic distributions for docs} \mid \alpha, \eta)$

New in version 2.0.0.

toLocal ()

Convert this distributed model to a local representation. This discards info about the training dataset.

WARNING: This involves collecting a large *topicsMatrix*() to the driver.

New in version 2.0.0.

trainingLogLikelihood ()

Log likelihood of the observed tokens in the training set, given the current parameter estimates: $\log P(\text{docs} \mid \text{topics, topic distributions for docs, Dirichlet hyperparameters})$

Notes:

- This excludes the prior; for that, use *logPrior*() .

- Even with `logPrior()`, this is NOT the same as the data log likelihood given the hyperparameters.
- This is computed from the topic distributions computed during training. If you call `logLikelihood()` on the same training dataset, the topic distributions will be computed again, possibly giving different results.

New in version 2.0.0.

```
class pyspark.ml.clustering.PowerIterationClustering (k=2, max-  
Iter=20,  
init-  
Mode='random',  
src-  
Col='src',  
dst-  
Col='dst',  
weight-  
Col=None)
```

Note: Experimental

Power Iteration Clustering (PIC), a scalable graph clustering algorithm developed by [Lin and Cohen](#). From the abstract: PIC finds a very low-dimensional embedding of a dataset using truncated power iteration on a normalized pair-wise similarity matrix of the data.

This class is not yet an Estimator/Transformer, use `assignClusters()` method to run the PowerIterationClustering algorithm.

See also:

[Wikipedia on Spectral clustering](#)

```
>>> data = [(1, 0, 0.5),  
...         (2, 0, 0.5), (2, 1, 0.7),  
...         (3, 0, 0.5), (3, 1, 0.7), (3, 2, 0.9),  
...         (4, 0, 0.5), (4, 1, 0.7), (4, 2, 0.9), (4, 3, 1.1),  
...         (5, 0, 0.5), (5, 1, 0.7), (5, 2, 0.9), (5, 3, 1.1), (5,  
↪ 4, 1.3)]  
>>> df = spark.createDataFrame(data).toDF("src", "dst", "weight")  
>>> pic = PowerIterationClustering(k=2, maxIter=40, weightCol=  
↪ "weight")  
>>> assignments = pic.assignClusters(df)  
>>> assignments.sort(assignments.id).show(truncate=False)  
+---+-----+  
|id |cluster|  
+---+-----+
```

(continues on next page)

(continued from previous page)

```

|0|1|
|1|1|
|2|1|
|3|1|
|4|1|
|5|0|
+---+-----+
...
>>> pic_path = temp_path + "/pic"
>>> pic.save(pic_path)
>>> pic2 = PowerIterationClustering.load(pic_path)
>>> pic2.getK()
2
>>> pic2.getMaxIter()
40

```

New in version 2.4.0.

assignClusters (*dataset*)

Run the PIC algorithm and returns a cluster assignment for each input vertex.

Parameters dataset – A dataset with columns src, dst, weight representing the affinity matrix, which is the matrix A in the PIC paper. Suppose the src column value is i , the dst column value is j , the weight column value is similarity s_{ij} , which must be nonnegative. This is a symmetric matrix and hence $s_{ij} = s_{ji}$. For any (i, j) with nonzero similarity, there should be either (i, j, s_{ij}) or (j, i, s_{ji}) in the input. Rows with $i = j$ are ignored, because we assume $s_{ij} = 0.0$.

Returns A dataset that contains columns of vertex id and the corresponding cluster for the id. The schema of it will be: - id: Long - cluster: Int

New in version 2.4.0.

New in version 2.4.0.

getDstCol ()

Gets the value of `dstCol` or its default value.

New in version 2.4.0.

getInitMode ()

Gets the value of `initMode` or its default value.

New in version 2.4.0.

getK ()

Gets the value of `k` or its default value.

New in version 2.4.0.

getSrcCol()

Gets the value of `srcCol` or its default value.

New in version 2.4.0.

setDstCol(value)

Sets the value of `dstCol`.

New in version 2.4.0.

setInitMode(value)

Sets the value of `initMode`.

New in version 2.4.0.

setK(value)

Sets the value of `k`.

New in version 2.4.0.

setParams(self, k=2, maxIter=20, initMode="random", srcCol="src", dstCol="dst", weightCol=None)

Sets params for `PowerIterationClustering`.

New in version 2.4.0.

setSrcCol(value)

Sets the value of `srcCol`.

New in version 2.4.0.

RECOMMENDATION API

```
class pyspark.ml.recommendation.ALS(rank=10,      maxIter=10,      reg-
                                     Param=0.1,      numUserBlocks=10,
                                     numItemBlocks=10,      implicit-
                                     Prefs=False,      alpha=1.0,      user-
                                     Col='user',      itemCol='item',
                                     seed=None,      ratingCol='rating',
                                     nonnegative=False,      checkpointIn-
                                     terval=10,      intermediateStor-
                                     ageLevel='MEMORY_AND_DISK',
                                     finalStor-
                                     ageLevel='MEMORY_AND_DISK',
                                     coldStartStrategy='nan')
```

Alternating Least Squares (ALS) matrix factorization.

ALS attempts to estimate the ratings matrix R as the product of two lower-rank matrices, X and Y , i.e. $X * Y^t = R$. Typically these approximations are called ‘factor’ matrices. The general approach is iterative. During each iteration, one of the factor matrices is held constant, while the other is solved for using least squares. The newly-solved factor matrix is then held constant while solving for the other factor matrix.

This is a blocked implementation of the ALS factorization algorithm that groups the two sets of factors (referred to as “users” and “products”) into blocks and reduces communication by only sending one copy of each user vector to each product block on each iteration, and only for the product blocks that need that user’s feature vector. This is achieved by pre-computing some information about the ratings matrix to determine the “out-links” of each user (which blocks of products it will contribute to) and “in-link” information for each product (which of the feature vectors it receives from each user block it will depend on). This allows us to send only an array of feature vectors between each user block and product block, and have the product block find the users’ ratings and update the products based on these messages.

For implicit preference data, the algorithm used is based on “[Collaborative Filtering for Implicit Feedback Datasets](#)”, adapted for the blocked approach used here.

Essentially instead of finding the low-rank approximations to the rating matrix R , this finds the approximations for a preference matrix P where the elements of P are 1 if $r > 0$ and 0

if $r \leq 0$. The ratings then act as ‘confidence’ values related to strength of indicated user preferences rather than explicit ratings given to items.

```
>>> df = spark.createDataFrame(
...     [(0, 0, 4.0), (0, 1, 2.0), (1, 1, 3.0), (1, 2, 4.0), (2, 1,
↪ 1.0), (2, 2, 5.0)],
...     ["user", "item", "rating"])
>>> als = ALS(rank=10, maxIter=5, seed=0)
>>> model = als.fit(df)
>>> model.rank
10
>>> model.userFactors.orderBy("id").collect()
[Row(id=0, features=[...]), Row(id=1, ...), Row(id=2, ...)]
>>> test = spark.createDataFrame([(0, 2), (1, 0), (2, 0)], ["user",
↪ "item"])
>>> predictions = sorted(model.transform(test).collect(), ↪
↪ key=lambda r: r[0])
>>> predictions[0]
Row(user=0, item=2, prediction=-0.13807615637779236)
>>> predictions[1]
Row(user=1, item=0, prediction=2.6258413791656494)
>>> predictions[2]
Row(user=2, item=0, prediction=-1.5018409490585327)
>>> user_recs = model.recommendForAllUsers(3)
>>> user_recs.where(user_recs.user == 0) .select(
↪ "recommendations.item", "recommendations.rating").collect()
[Row(item=[0, 1, 2], rating=[3.910..., 1.992..., -0.138...])]
>>> item_recs = model.recommendForAllItems(3)
>>> item_recs.where(item_recs.item == 2) .select(
↪ "recommendations.user", "recommendations.rating").collect()
[Row(user=[2, 1, 0], rating=[4.901..., 3.981..., -0.138...])]
>>> user_subset = df.where(df.user == 2)
>>> user_subset_recs = model.recommendForUserSubset(user_subset, 3)
>>> user_subset_recs.select("recommendations.item",
↪ "recommendations.rating").first()
Row(item=[2, 1, 0], rating=[4.901..., 1.056..., -1.501...])
>>> item_subset = df.where(df.item == 0)
>>> item_subset_recs = model.recommendForItemSubset(item_subset, 3)
>>> item_subset_recs.select("recommendations.user",
↪ "recommendations.rating").first()
Row(user=[0, 1, 2], rating=[3.910..., 2.625..., -1.501...])
>>> als_path = temp_path + "/als"
>>> als.save(als_path)
>>> als2 = ALS.load(als_path)
>>> als.getMaxIter()
5
>>> model_path = temp_path + "/als_model"
```

(continues on next page)

(continued from previous page)

```

>>> model.save(model_path)
>>> model2 = ALSModel.load(model_path)
>>> model.rank == model2.rank
True
>>> sorted(model.userFactors.collect()) == sorted(model2.
    ↳userFactors.collect())
True
>>> sorted(model.itemFactors.collect()) == sorted(model2.
    ↳itemFactors.collect())
True

```

New in version 1.4.0.

getAlpha()

Gets the value of alpha or its default value.

New in version 1.4.0.

getColdStartStrategy()

Gets the value of coldStartStrategy or its default value.

New in version 2.2.0.

getFinalStorageLevel()

Gets the value of finalStorageLevel or its default value.

New in version 2.0.0.

getImplicitPrefs()

Gets the value of implicitPrefs or its default value.

New in version 1.4.0.

getIntermediateStorageLevel()

Gets the value of intermediateStorageLevel or its default value.

New in version 2.0.0.

getItemCol()

Gets the value of itemCol or its default value.

New in version 1.4.0.

getNonnegative()

Gets the value of nonnegative or its default value.

New in version 1.4.0.

getNumItemBlocks()

Gets the value of numItemBlocks or its default value.

New in version 1.4.0.

getNumUserBlocks ()

Gets the value of numUserBlocks or its default value.

New in version 1.4.0.

getRank ()

Gets the value of rank or its default value.

New in version 1.4.0.

getRatingCol ()

Gets the value of ratingCol or its default value.

New in version 1.4.0.

getUserCol ()

Gets the value of userCol or its default value.

New in version 1.4.0.

setAlpha (value)

Sets the value of alpha.

New in version 1.4.0.

setColdStartStrategy (value)

Sets the value of coldStartStrategy.

New in version 2.2.0.

setFinalStorageLevel (value)

Sets the value of finalStorageLevel.

New in version 2.0.0.

setImplicitPrefs (value)

Sets the value of implicitPrefs.

New in version 1.4.0.

setIntermediateStorageLevel (value)

Sets the value of intermediateStorageLevel.

New in version 2.0.0.

setItemCol (value)

Sets the value of itemCol.

New in version 1.4.0.

setNonnegative (value)

Sets the value of nonnegative.

New in version 1.4.0.

setNumBlocks (*value*)

Sets both `numUserBlocks` and `numItemBlocks` to the specific value.

New in version 1.4.0.

setNumItemBlocks (*value*)

Sets the value of `numItemBlocks`.

New in version 1.4.0.

setNumUserBlocks (*value*)

Sets the value of `numUserBlocks`.

New in version 1.4.0.

setParams (*self*, *rank*=10, *maxIter*=10, *regParam*=0.1, *numUserBlocks*=10, *numItemBlocks*=10, *implicitPrefs*=False, *alpha*=1.0, *userCol*="user", *itemCol*="item", *seed*=None, *ratingCol*="rating", *nonnegative*=False, *checkpointInterval*=10, *intermediateStorageLevel*="MEMORY_AND_DISK", *finalStorageLevel*="MEMORY_AND_DISK", *coldStartStrategy*="nan")

Sets params for ALS.

New in version 1.4.0.

setRank (*value*)

Sets the value of `rank`.

New in version 1.4.0.

setRatingCol (*value*)

Sets the value of `ratingCol`.

New in version 1.4.0.

setUserCol (*value*)

Sets the value of `userCol`.

New in version 1.4.0.

class pyspark.ml.recommendation.**ALSModel** (*java_model*=None)

Model fitted by ALS.

New in version 1.4.0.

itemFactors

id and *features*

New in version 1.4.0.

Type a DataFrame that stores item factors in two columns

rank

rank of the matrix factorization model

New in version 1.4.0.

recommendForAllItems (*numUsers*)

Returns top *numUsers* users recommended for each item, for all items.

Parameters **numUsers** – max number of recommendations for each item

Returns a DataFrame of (itemCol, recommendations), where recommendations are stored as an array of (userCol, rating) Rows.

New in version 2.2.0.

recommendForAllUsers (*numItems*)

Returns top *numItems* items recommended for each user, for all users.

Parameters **numItems** – max number of recommendations for each user

Returns a DataFrame of (userCol, recommendations), where recommendations are stored as an array of (itemCol, rating) Rows.

New in version 2.2.0.

recommendForItemSubset (*dataset*, *numUsers*)

Returns top *numUsers* users recommended for each item id in the input data set. Note that if there are duplicate ids in the input dataset, only one set of recommendations per unique id will be returned.

Parameters

- **dataset** – a Dataset containing a column of item ids. The column name must match *itemCol*.
- **numUsers** – max number of recommendations for each item

Returns a DataFrame of (itemCol, recommendations), where recommendations are stored as an array of (userCol, rating) Rows.

New in version 2.3.0.

recommendForUserSubset (*dataset*, *numItems*)

Returns top *numItems* items recommended for each user id in the input data set. Note that if there are duplicate ids in the input dataset, only one set of recommendations per unique id will be returned.

Parameters

- **dataset** – a Dataset containing a column of user ids. The column name must match *userCol*.
- **numItems** – max number of recommendations for each user

Returns a DataFrame of (userCol, recommendations), where recommendations are stored as an array of (itemCol, rating) Rows.

New in version 2.3.0.

userFactors

id and features

New in version 1.4.0.

Type a DataFrame that stores user factors in two columns

PIPELINE API

class `pyspark.ml.pipeline.Pipeline` (*stages=None*)

A simple pipeline, which acts as an estimator. A Pipeline consists of a sequence of stages, each of which is either an `Estimator` or a `Transformer`. When `Pipeline.fit()` is called, the stages are executed in order. If a stage is an `Estimator`, its `Estimator.fit()` method will be called on the input dataset to fit a model. Then the model, which is a transformer, will be used to transform the dataset as the input to the next stage. If a stage is a `Transformer`, its `Transformer.transform()` method will be called to produce the dataset for the next stage. The fitted model from a *Pipeline* is a *PipelineModel*, which consists of fitted models and transformers, corresponding to the pipeline stages. If stages is an empty list, the pipeline acts as an identity transformer.

New in version 1.3.0.

copy (*extra=None*)

Creates a copy of this instance.

Parameters *extra* – extra parameters

Returns new instance

New in version 1.4.0.

getStages ()

Get pipeline stages.

New in version 1.3.0.

classmethod read ()

Returns an `MLReader` instance for this class.

New in version 2.0.0.

setParams (*self*, *stages=None*)

Sets params for Pipeline.

New in version 1.3.0.

setStages (*value*)

Set pipeline stages.

Parameters **value** – a list of transformers or estimators

Returns the pipeline instance

New in version 1.3.0.

write()

Returns an MLWriter instance for this ML instance.

New in version 2.0.0.

class pyspark.ml.pipeline.**PipelineModel**(*stages*)

Represents a compiled pipeline with transformers and fitted models.

New in version 1.3.0.

copy(*extra=None*)

Creates a copy of this instance.

Parameters **extra** – extra parameters

Returns new instance

New in version 1.4.0.

classmethod **read()**

Returns an MLReader instance for this class.

New in version 2.0.0.

write()

Returns an MLWriter instance for this ML instance.

New in version 2.0.0.

class pyspark.ml.pipeline.**PipelineModelReader**(*cls*)

(Private) Specialization of MLReader for *PipelineModel* types

load(*path*)

Load the ML instance from the input path.

class pyspark.ml.pipeline.**PipelineModelWriter**(*instance*)

(Private) Specialization of MLWriter for *PipelineModel* types

saveImpl(*path*)

save() handles overwriting and then calls this method. Subclasses should override this method to implement the actual saving of the instance.

class pyspark.ml.pipeline.**PipelineReader**(*cls*)

(Private) Specialization of MLReader for *Pipeline* types

load(*path*)

Load the ML instance from the input path.

```
class pyspark.ml.pipeline.PipelineSharedReadWrite
```

Note: DeveloperApi

Functions for MLReader and MLWriter shared between *Pipeline* and *PipelineModel*

New in version 2.3.0.

static **getStagePath** (*stageUid*, *stageIdx*, *numStages*, *stagesDir*)

Get path for saving the given stage.

static **load** (*metadata*, *sc*, *path*)

Load metadata and stages for a *Pipeline* or *PipelineModel*

Returns (UID, list of stages)

static **saveImpl** (*instance*, *stages*, *sc*, *path*)

Save metadata and stages for a *Pipeline* or *PipelineModel* - save metadata to path/metadata - save stages to stages/IDX_UID

static **validateStages** (*stages*)

Check that all stages are Writable

```
class pyspark.ml.pipeline.PipelineWriter (instance)
```

(Private) Specialization of MLWriter for *Pipeline* types

saveImpl (*path*)

save() handles overwriting and then calls this method. Subclasses should override this method to implement the actual saving of the instance.

TUNING API

class pyspark.ml.tuning.ParamGridBuilder

Builder for a param grid used in grid search-based model selection.

```
>>> from pyspark.ml.classification import LogisticRegression
>>> lr = LogisticRegression()
>>> output = ParamGridBuilder() \
...     .baseOn({lr.labelCol: 'l'}) \
...     .baseOn([lr.predictionCol, 'p']) \
...     .addGrid(lr.regParam, [1.0, 2.0]) \
...     .addGrid(lr.maxIter, [1, 5]) \
...     .build()
>>> expected = [
...     {lr.regParam: 1.0, lr.maxIter: 1, lr.labelCol: 'l', lr.
... ↪predictionCol: 'p'},
...     {lr.regParam: 2.0, lr.maxIter: 1, lr.labelCol: 'l', lr.
... ↪predictionCol: 'p'},
...     {lr.regParam: 1.0, lr.maxIter: 5, lr.labelCol: 'l', lr.
... ↪predictionCol: 'p'},
...     {lr.regParam: 2.0, lr.maxIter: 5, lr.labelCol: 'l', lr.
... ↪predictionCol: 'p'}]
>>> len(output) == len(expected)
True
>>> all([m in expected for m in output])
True
```

New in version 1.4.0.

addGrid (*param, values*)

Sets the given parameters in this grid to fixed values.

New in version 1.4.0.

baseOn (**args*)

Sets the given parameters in this grid to fixed values. Accepts either a parameter dictionary or a list of (parameter, value) pairs.

New in version 1.4.0.

build()

Builds and returns all combinations of parameters specified by the param grid.

New in version 1.4.0.

class pyspark.ml.tuning.**CrossValidator** (*estimator=None, estimator-ParamMaps=None, evaluator=None, numFolds=3, seed=None, parallelism=1, collectSubModels=False*)

K-fold cross validation performs model selection by splitting the dataset into a set of non-overlapping randomly partitioned folds which are used as separate training and test datasets e.g., with k=3 folds, K-fold cross validation will generate 3 (training, test) dataset pairs, each of which uses 2/3 of the data for training and 1/3 for testing. Each fold is used as the test set exactly once.

```
>>> from pyspark.ml.classification import LogisticRegression
>>> from pyspark.ml.evaluation import BinaryClassificationEvaluator
>>> from pyspark.ml.linalg import Vectors
>>> dataset = spark.createDataFrame(
...     [(Vectors.dense([0.0]), 0.0),
...      (Vectors.dense([0.4]), 1.0),
...      (Vectors.dense([0.5]), 0.0),
...      (Vectors.dense([0.6]), 1.0),
...      (Vectors.dense([1.0]), 1.0)] * 10,
...     ["features", "label"])
>>> lr = LogisticRegression()
>>> grid = ParamGridBuilder().addGrid(lr.maxIter, [0, 1]).build()
>>> evaluator = BinaryClassificationEvaluator()
>>> cv = CrossValidator(estimator=lr, estimatorParamMaps=grid,
→evaluator=evaluator,
...     parallelism=2)
>>> cvModel = cv.fit(dataset)
>>> cvModel.avgMetrics[0]
0.5
>>> evaluator.evaluate(cvModel.transform(dataset))
0.8333...
```

New in version 1.4.0.

copy (*extra=None*)

Creates a copy of this instance with a randomly generated uid and some extra params. This copies creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over.

Parameters **extra** – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 1.4.0.

getNumFolds()

Gets the value of numFolds or its default value.

New in version 1.4.0.

classmethod read()

Returns an MLReader instance for this class.

New in version 2.3.0.

setNumFolds(value)

Sets the value of numFolds.

New in version 1.4.0.

setParams(estimator=None, estimatorParamMaps=None, evaluator=None, numFolds=3, seed=None, parallelism=1, collectSubModels=False)

setParams(self, estimator=None, estimatorParamMaps=None, evaluator=None, numFolds=3, seed=None, parallelism=1, collectSubModels=False): Sets params for cross validator.

New in version 1.4.0.

write()

Returns an MLWriter instance for this ML instance.

New in version 2.3.0.

class pyspark.ml.tuning.**CrossValidatorModel**(*bestModel*, *avgMetrics*,
rics=[], *subModels=None*)

CrossValidatorModel contains the model with the highest average cross-validation metric across folds and uses this model to transform input data. CrossValidatorModel also tracks the metrics for each param map evaluated.

New in version 1.4.0.

avgMetrics = None

Average cross-validation metrics for each paramMap in CrossValidator.estimatorParamMaps, in the corresponding order.

bestModel = None

best model from cross validation

copy(extra=None)

Creates a copy of this instance with a randomly generated uid and some extra params. This copies the underlying bestModel, creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over. It does not copy the extra Params into the subModels.

Parameters **extra** – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 1.4.0.

classmethod read()

Returns an MLReader instance for this class.

New in version 2.3.0.

subModels = None

sub model list from cross validation

write()

Returns an MLWriter instance for this ML instance.

New in version 2.3.0.

```
class pyspark.ml.tuning.TrainValidationSplit (estimator=None,  
estimator-  
ParamMaps=None,  
evaluator=None,  
trainRatio=0.75,  
parallelism=1, col-  
lectSubModels=False,  
seed=None)
```

Note: Experimental

Validation for hyper-parameter tuning. Randomly splits the input dataset into train and validation sets, and uses evaluation metric on the validation set to select the best model. Similar to *CrossValidator*, but only splits the set once.

```
>>> from pyspark.ml.classification import LogisticRegression
>>> from pyspark.ml.evaluation import BinaryClassificationEvaluator
>>> from pyspark.ml.linalg import Vectors
>>> dataset = spark.createDataFrame(
...     [(Vectors.dense([0.0]), 0.0),
...      (Vectors.dense([0.4]), 1.0),
...      (Vectors.dense([0.5]), 0.0),
...      (Vectors.dense([0.6]), 1.0),
...      (Vectors.dense([1.0]), 1.0)] * 10,
...     ["features", "label"])
>>> lr = LogisticRegression()
>>> grid = ParamGridBuilder().addGrid(lr.maxIter, [0, 1]).build()
>>> evaluator = BinaryClassificationEvaluator()
>>> tvs = TrainValidationSplit(estimator=lr,
→estimatorParamMaps=grid, evaluator=evaluator, (continues on next page)
```


(continued from previous page)

```

...     parallelism=2)
>>> tvsModel = tvs.fit(dataset)
>>> evaluator.evaluate(tvsModel.transform(dataset))
0.8333...

```

New in version 2.0.0.

copy (*extra=None*)

Creates a copy of this instance with a randomly generated uid and some extra params. This copies creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over.

Parameters *extra* – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 2.0.0.

getTrainRatio ()

Gets the value of trainRatio or its default value.

New in version 2.0.0.

classmethod read ()

Returns an MLReader instance for this class.

New in version 2.3.0.

setParams (*estimator=None, estimatorParamMaps=None, evaluator=None, train-Ratio=0.75, parallelism=1, collectSubModels=False, seed=None*)

setParams(self, estimator=None, estimatorParamMaps=None, evaluator=None, train-Ratio=0.75, parallelism=1, collectSubModels=False, seed=None): Sets params for the train validation split.

New in version 2.0.0.

setTrainRatio (*value*)

Sets the value of trainRatio.

New in version 2.0.0.

write ()

Returns an MLWriter instance for this ML instance.

New in version 2.3.0.

class pyspark.ml.tuning.**TrainValidationSplitModel** (*bestModel, validation-Metrics=[], sub-Models=None*)

Note: Experimental

Model from train validation split.

New in version 2.0.0.

bestModel = None

best model from train validation split

copy (*extra=None*)

Creates a copy of this instance with a randomly generated uid and some extra params. This copies the underlying bestModel, creates a deep copy of the embedded paramMap, and copies the embedded and extra parameters over. And, this creates a shallow copy of the validationMetrics. It does not copy the extra Params into the subModels.

Parameters **extra** – Extra parameters to copy to the new instance

Returns Copy of this instance

New in version 2.0.0.

classmethod read ()

Returns an MLReader instance for this class.

New in version 2.3.0.

subModels = None

sub models from train validation split

validationMetrics = None

evaluated validation metrics

write ()

Returns an MLWriter instance for this ML instance.

New in version 2.3.0.

EVALUATION API

class pyspark.ml.evaluation.Evaluator

Base class for evaluators that compute metrics from predictions.

New in version 1.4.0.

evaluate (*dataset*, *params=None*)

Evaluates the output with optional parameters.

Parameters

- **dataset** – a dataset that contains labels/observations and predictions
- **params** – an optional param map that overrides embedded params

Returns

metric

New in version 1.4.0.

isLargerBetter ()

Indicates whether the metric returned by *evaluate()* should be maximized (True, default) or minimized (False). A given evaluator may support multiple metrics which may be maximized or minimized.

New in version 1.5.0.

class pyspark.ml.evaluation.BinaryClassificationEvaluator (*rawPredictionCol='rawP*
la-
bel-
Col='label',
met-
ric-
Name='areaUnderROC')

Note: Experimental

Evaluator for binary classification, which expects two input columns: `rawPrediction` and `label`. The `rawPrediction` column can be of type double (binary 0/1 prediction, or probability of label 1) or of type vector (length-2 vector of raw predictions, scores, or label probabilities).

```
>>> from pyspark.ml.linalg import Vectors
>>> scoreAndLabels = map(lambda x: (Vectors.dense([1.0 - x[0],
↳ x[0]]), x[1]),
... [(0.1, 0.0), (0.1, 1.0), (0.4, 0.0), (0.6, 0.0), (0.6, 1.0),
↳ (0.6, 1.0), (0.8, 1.0)])
>>> dataset = spark.createDataFrame(scoreAndLabels, ["raw", "label"
↳ ])
...
>>> evaluator = BinaryClassificationEvaluator(rawPredictionCol="raw"
↳ )
>>> evaluator.evaluate(dataset)
0.70...
>>> evaluator.evaluate(dataset, {evaluator.metricName: "areaUnderPR"
↳ })
0.83...
>>> bce_path = temp_path + "/bce"
>>> evaluator.save(bce_path)
>>> evaluator2 = BinaryClassificationEvaluator.load(bce_path)
>>> str(evaluator2.getRawPredictionCol())
'raw'
```

New in version 1.4.0.

getMetricName()

Gets the value of `metricName` or its default value.

New in version 1.4.0.

setMetricName(value)

Sets the value of `metricName`.

New in version 1.4.0.

**setParams(self, rawPredictionCol="rawPrediction", labelCol="label", metric-
Name="areaUnderROC")**

Sets params for binary classification evaluator.

New in version 1.4.0.

```
class pyspark.ml.evaluation.RegressionEvaluator (predictionCol='prediction',
labelCol='label',
metric-
Name='rmse')
```

Note: Experimental

Evaluator for Regression, which expects two input columns: prediction and label.

```
>>> scoreAndLabels = [(-28.98343821, -27.0), (20.21491975, 21.5),
...                    (-25.98418959, -22.0), (30.69731842, 33.0), (74.69283752, 71.
    ↪0)]
>>> dataset = spark.createDataFrame(scoreAndLabels, ["raw", "label
    ↪"])
...
>>> evaluator = RegressionEvaluator(predictionCol="raw")
>>> evaluator.evaluate(dataset)
2.842...
>>> evaluator.evaluate(dataset, {evaluator.metricName: "r2"})
0.993...
>>> evaluator.evaluate(dataset, {evaluator.metricName: "mae"})
2.649...
>>> re_path = temp_path + "/re"
>>> evaluator.save(re_path)
>>> evaluator2 = RegressionEvaluator.load(re_path)
>>> str(evaluator2.getPredictionCol())
'raw'
```

New in version 1.4.0.

getMetricName()

Gets the value of metricName or its default value.

New in version 1.4.0.

setMetricName(value)

Sets the value of metricName.

New in version 1.4.0.

setParams(self, predictionCol="prediction", labelCol="label", metric-
Name="rmse")

Sets params for regression evaluator.

New in version 1.4.0.

class pyspark.ml.evaluation.**MulticlassClassificationEvaluator** (predictionCol='pre-
 la-
 bel-
 Col='label',
 met-
 ric-
 Name='f1')

Note: Experimental

Evaluator for Multiclass Classification, which expects two input columns: prediction and label.

```
>>> scoreAndLabels = [(0.0, 0.0), (0.0, 1.0), (0.0, 0.0),
...                    (1.0, 0.0), (1.0, 1.0), (1.0, 1.0), (1.0, 1.0), (2.0, 2.0),
↪ (2.0, 0.0)]
>>> dataset = spark.createDataFrame(scoreAndLabels, ["prediction",
↪ "label"])
...
>>> evaluator = MulticlassClassificationEvaluator(predictionCol=
↪ "prediction")
>>> evaluator.evaluate(dataset)
0.66...
>>> evaluator.evaluate(dataset, {evaluator.metricName: "accuracy"})
0.66...
>>> mce_path = temp_path + "/mce"
>>> evaluator.save(mce_path)
>>> evaluator2 = MulticlassClassificationEvaluator.load(mce_path)
>>> str(evaluator2.getPredictionCol())
'prediction'
```

New in version 1.5.0.

getMetricName()

Gets the value of metricName or its default value.

New in version 1.5.0.

setMetricName(value)

Sets the value of metricName.

New in version 1.5.0.

setParams(self, predictionCol="prediction", labelCol="label", metric-
Name="f1")

Sets params for multiclass classification evaluator.

New in version 1.5.0.

class pyspark.ml.evaluation.**ClusteringEvaluator** (*predictionCol='prediction',*
fea-
turesCol='features',
metric-
Name='silhouette',
distanceMea-
sure='squaredEuclidean')

Note: Experimental

Evaluator for Clustering results, which expects two input columns: prediction and features. The metric computes the Silhouette measure using the squared Euclidean distance.

The Silhouette is a measure for the validation of the consistency within clusters. It ranges between 1 and -1, where a value close to 1 means that the points in a cluster are close to the other points in the same cluster and far from the points of the other clusters.

```
>>> from pyspark.ml.linalg import Vectors
>>> featureAndPredictions = map(lambda x: (Vectors.dense(x[0]),
    ↪x[1]),
...     [[(0.0, 0.5], 0.0), ([0.5, 0.0], 0.0), ([10.0, 11.0], 1.0),
...       ([10.5, 11.5], 1.0), ([1.0, 1.0], 0.0), ([8.0, 6.0], 1.0)])
>>> dataset = spark.createDataFrame(featureAndPredictions, [
    ↪"features", "prediction"])
...
>>> evaluator = ClusteringEvaluator(predictionCol="prediction")
>>> evaluator.evaluate(dataset)
0.9079...
>>> ce_path = temp_path + "/ce"
>>> evaluator.save(ce_path)
>>> evaluator2 = ClusteringEvaluator.load(ce_path)
>>> str(evaluator2.getPredictionCol())
'prediction'
```

New in version 2.3.0.

getDistanceMeasure()

Gets the value of *distanceMeasure*

New in version 2.4.0.

getMetricName()

Gets the value of *metricName* or its default value.

New in version 2.3.0.

setDistanceMeasure(value)

Sets the value of *distanceMeasure*.

New in version 2.4.0.

setMetricName(value)

Sets the value of *metricName*.

New in version 2.3.0.

```
setParams (self, predictionCol="prediction", featuresCol="features", metric-  
             Name="silhouette", distanceMeasure="squaredEuclidean")
```

Sets params for clustering evaluator.

New in version 2.3.0.

CHAPTER

TEN

MAIN REFERENCE

BIBLIOGRAPHY

[PySpark] [Apache Spark](#).

PYTHON MODULE INDEX

p

- `pyspark.ml.classification`, [39](#)
- `pyspark.ml.clustering`, [67](#)
- `pyspark.ml.evaluation`, [103](#)
- `pyspark.ml.pipeline`, [93](#)
- `pyspark.ml.recommendation`, [85](#)
- `pyspark.ml.regression`, [13](#)
- `pyspark.ml.stat`, [5](#)
- `pyspark.ml.tuning`, [97](#)

INDEX

A

- accuracy (pyspark.ml.classification.LogisticRegressionSummary attribute), 46
- addGrid() (pyspark.ml.tuning.ParamGridBuilder method), 97
- AFTSurvivalRegression (class in pyspark.ml.regression), 13
- AFTSurvivalRegressionModel (class in pyspark.ml.regression), 15
- aic (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 24
- ALS (class in pyspark.ml.recommendation), 85
- ALSModel (class in pyspark.ml.recommendation), 89
- areaUnderROC (pyspark.ml.classification.BinaryLogisticRegressionSummary attribute), 48
- assignClusters() (pyspark.ml.clustering.PowerIterationClustering method), 83
- avgMetrics (pyspark.ml.tuning.CrossValidatorModel attribute), 99

B

- baseOn() (pyspark.ml.tuning.ParamGridBuilder method), 97
- bestModel (pyspark.ml.tuning.CrossValidatorModel attribute), 99
- bestModel (pyspark.ml.tuning.TrainValidationSplitModel

attribute), 102

- BinaryClassificationEvaluator (class in pyspark.ml.evaluation), 103
- BinaryLogisticRegressionSummary (class in pyspark.ml.classification), 48
- BinaryLogisticRegressionTrainingSummary (class in pyspark.ml.classification), 49
- BisectingKMeans (class in pyspark.ml.clustering), 67
- BisectingKMeansModel (class in pyspark.ml.clustering), 69
- BisectingKMeansSummary (class in pyspark.ml.clustering), 69
- boundaries (pyspark.ml.regression.IsotonicRegressionModel attribute), 28
- build() (pyspark.ml.tuning.ParamGridBuilder method), 98

C

- ChiSquareTest (class in pyspark.ml.stat), 5
- clusterCenters() (pyspark.ml.clustering.BisectingKMeansModel method), 69
- clusterCenters() (pyspark.ml.clustering.KMeansModel method), 72
- ClusteringEvaluator (class in pyspark.ml.evaluation), 106
- coefficientMatrix (pyspark.ml.classification.LogisticRegressionModel attribute), 45
- coefficients (pyspark.ml.classification.LinearSVCModel

attribute), 40
 coefficients (pyspark.ml.classification.LogisticRegressionModel attribute), 45
 coefficients (pyspark.ml.regression.AFTSurvivalRegressionModel attribute), 15
 coefficients (pyspark.ml.regression.GeneralizedLinearRegressionModel attribute), 24
 coefficients (pyspark.ml.regression.LinearRegressionModel attribute), 30
 coefficientStandardErrors (pyspark.ml.regression.GeneralizedLinearRegressionModel attribute), 26
 coefficientStandardErrors (pyspark.ml.regression.LinearRegressionSummary attribute), 31
 computeCost() (pyspark.ml.clustering.BisectingKMeansModel method), 69
 computeCost() (pyspark.ml.clustering.KMeansModel method), 72
 copy() (pyspark.ml.classification.OneVsRestModel method), 65
 copy() (pyspark.ml.classification.OneVsRestModel method), 65
 copy() (pyspark.ml.pipeline.Pipeline method), 93
 copy() (pyspark.ml.pipeline.PipelineModel method), 94
 copy() (pyspark.ml.tuning.CrossValidator method), 98
 copy() (pyspark.ml.tuning.CrossValidatorModel method), 99
 copy() (pyspark.ml.tuning.TrainValidationSplit method), 101
 copy() (pyspark.ml.tuning.TrainValidationSplitModel method), 102
 corr() (pyspark.ml.stat.Correlation static method), 6
 correlation (class in pyspark.ml.stat), 6
 count() (pyspark.ml.stat.Summarizer static method), 9
 CrossValidator (class in pyspark.ml.tuning), 98
 CrossValidatorModel (class in pyspark.ml.tuning), 99
D
 DecisionTreeClassificationModel (class in pyspark.ml.classification), 52
 DecisionTreeClassifier (class in pyspark.ml.classification), 50
 DecisionTreeRegressionModel (class in pyspark.ml.regression), 17
 DecisionTreeRegressor (class in pyspark.ml.regression), 16
 degreesOfFreedom (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 degreesOfFreedom (pyspark.ml.regression.LinearRegressionSummary attribute), 31
 describeTopics() (pyspark.ml.clustering.LDAModel method), 79
 deviance (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 devianceResiduals (pyspark.ml.regression.LinearRegressionSummary attribute), 31
 dispersion (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 DistributedLDAModel (class in pyspark.ml.clustering), 81
E
 estimatedDocConcentration() (pyspark.ml.clustering.LDAModel method), 80

[evaluate\(\)](#) (pyspark.ml.classification.LogisticRegressionModel attribute), 36
[evaluate\(\)](#) (pyspark.ml.classification.LogisticRegressionSummary method), 45
[evaluate\(\)](#) (pyspark.ml.evaluation.Evaluator method), 103
[evaluate\(\)](#) (pyspark.ml.regression.LinearRegressionSummary attribute), 32
[evaluate\(\)](#) (pyspark.ml.regression.GeneralizedLinearRegressionModel method), 24
[evaluate\(\)](#) (pyspark.ml.regression.LinearRegressionModel method), 30
[evaluateEachIteration\(\)](#) (pyspark.ml.classification.GBTClassificationModel method), 55
[evaluateEachIteration\(\)](#) (pyspark.ml.regression.GBTRegressionModel method), 20
[Evaluator](#) (class in pyspark.ml.evaluation), 103
[explainedVariance](#) (pyspark.ml.regression.LinearRegressionSummary attribute), 31
F
[falsePositiveRateByLabel](#) (pyspark.ml.classification.LogisticRegressionSummary attribute), 46
[featureImportances](#) (pyspark.ml.classification.DecisionTreeClassificationModel attribute), 52
[featureImportances](#) (pyspark.ml.classification.GBTClassificationModel attribute), 55
[featureImportances](#) (pyspark.ml.classification.RandomForestClassificationModel attribute), 58
[featureImportances](#) (pyspark.ml.regression.DecisionTreeRegressionModel attribute), 17
[featureImportances](#) (pyspark.ml.regression.GBTRegressionModel attribute), 20
[featureImportances](#) (pyspark.ml.regression.RandomForestRegressionModel attribute), 20
[featuresCol](#) (pyspark.ml.classification.LogisticRegressionSummary attribute), 46
[featuresCol](#) (pyspark.ml.regression.LinearRegressionSummary attribute), 32
[fit](#) (pyspark.ml.classification.LogisticRegressionSummary method), 46
[fit](#) (pyspark.ml.classification.BinaryLogisticRegressionSummary method), 48
G
[GaussianMixture](#) (class in pyspark.ml.clustering), 72
[GaussianMixtureModel](#) (class in pyspark.ml.clustering), 74
[GaussianMixtureSummary](#) (class in pyspark.ml.clustering), 75
[gaussiansDF](#) (pyspark.ml.clustering.GaussianMixtureModel attribute), 74
[GBTClassificationModel](#) (class in pyspark.ml.classification), 55
[GBTClassifier](#) (class in pyspark.ml.classification), 52
[GBTRegressionModel](#) (class in pyspark.ml.regression), 20
[GBRegressor](#) (class in pyspark.ml.regression), 18
[GeneralizedLinearRegression](#) (class in pyspark.ml.regression), 20
[GeneralizedLinearRegressionModel](#) (class in pyspark.ml.regression), 24
[GeneralizedLinearRegressionSummary](#) (class in pyspark.ml.regression), 24
[GeneralizedLinearRegressionTrainingSummary](#) (class in pyspark.ml.regression), 26
[getAlpha\(\)](#) (pyspark.ml.recommendation.ALS method), 87
[getBlockSize\(\)](#) (pyspark.ml.classification.MultilayerPerceptronClassifier

<i>method</i>), 62	<i>method</i>), 62
getCensorCol() (pyspark.ml.regression.AFTSurvivalRegression <i>method</i>), 14	getInitMode() (pyspark.ml.clustering.KMeans <i>method</i>), 71
getCheckpointFiles() (pyspark.ml.clustering.DistributedLDAModel <i>method</i>), 81	getInitMode() (pyspark.ml.clustering.PowerIterationClustering <i>method</i>), 83
getColdStartStrategy() (pyspark.ml.recommendation.ALS <i>method</i>), 87	getInitSteps() (pyspark.ml.clustering.KMeans <i>method</i>), 71
getDistanceMeasure() (pyspark.ml.clustering.BisectingKMeans <i>method</i>), 68	getIntermediateStorageLevel() (pyspark.ml.recommendation.ALS <i>method</i>), 87
getDistanceMeasure() (pyspark.ml.clustering.KMeans <i>method</i>), 71	getIsotonic() (pyspark.ml.regression.IsotonicRegression <i>method</i>), 27
getDistanceMeasure() (pyspark.ml.evaluation.ClusteringEvaluator <i>method</i>), 107	getItemCol() (pyspark.ml.recommendation.ALS <i>method</i>), 87
getDocConcentration() (pyspark.ml.clustering.LDA <i>method</i>), 76	getK() (pyspark.ml.clustering.BisectingKMeans <i>method</i>), 68
getDstCol() (pyspark.ml.clustering.PowerIterationClustering <i>method</i>), 83	getK() (pyspark.ml.clustering.GaussianMixture <i>method</i>), 74
getEpsilon() (pyspark.ml.regression.LinearRegression <i>method</i>), 30	getK() (pyspark.ml.clustering.KMeans <i>method</i>), 71
getFamily() (pyspark.ml.classification.LogisticRegression <i>method</i>), 43	getK() (pyspark.ml.clustering.LDA <i>method</i>), 76
getFamily() (pyspark.ml.regression.GeneralizedLinearRegression <i>method</i>), 22	getK() (pyspark.ml.clustering.PowerIterationClustering <i>method</i>), 83
getFeatureIndex() (pyspark.ml.regression.IsotonicRegression <i>method</i>), 27	getLastCheckpoint() (pyspark.ml.clustering.LDA <i>method</i>), 77
getFinalStorageLevel() (pyspark.ml.recommendation.ALS <i>method</i>), 87	getLayers() (pyspark.ml.classification.MultilayerPerceptronClassifier <i>method</i>), 63
getImplicitPrefs() (pyspark.ml.recommendation.ALS <i>method</i>), 87	getLearningDecay() (pyspark.ml.clustering.LDA <i>method</i>), 77
getInitialWeights() (pyspark.ml.classification.MultilayerPerceptronClassifier <i>method</i>), 63	getLearningOffset() (pyspark.ml.clustering.LDA <i>method</i>), 77
	getMaxCheckpoint() (pyspark.ml.clustering.LDA <i>method</i>), 77

<i>park.ml.regression.GeneralizedLinearRegression</i> <i>park.ml.recommendation.ALS</i> method), method), 22	87
getLinkPower() (pys- <i>park.ml.regression.GeneralizedLinearRegression</i> <i>park.ml.recommendation.ALS</i> method), method), 23	87
getLinkPredictionCol() (pys- <i>park.ml.regression.GeneralizedLinearRegression</i> <i>park.ml.regression.GeneralizedLinearRegression</i> method), 23	method), 23
getLossType() (pys- <i>park.ml.classification.GBTClassifier</i> <i>park.ml.classification.GBTClassifier</i> method), 54	getOptimizeDocConcentration() (pyspark.ml.clustering.LDA method), 77
getLossType() (pys- <i>park.ml.regression.GBTRegressor</i> <i>park.ml.regression.GBTRegressor</i> method), 19	getOptimizer() (pyspark.ml.clustering.LDA method), 77
getLowerBoundsOnCoefficients() (pyspark.ml.classification.LogisticRegression method), 43	getQuantileProbabilities() (pyspark.ml.regression.AFTSurvivalRegression method), 14
getLowerBoundsOnIntercepts() (pyspark.ml.classification.LogisticRegression method), 43	getQuantilesCol() (pyspark.ml.regression.AFTSurvivalRegression method), 14
getMetricName() (pyspark.ml.evaluation.BinaryClassificationEvaluator <i>park.ml.recommendation.ALS</i> method), method), 104	getRank() (pyspark.ml.recommendation.ALS method), 88
getMetricName() (pyspark.ml.evaluation.ClusteringEvaluator <i>park.ml.recommendation.ALS</i> method), method), 107	getRatingCol() (pyspark.ml.recommendation.ALS method), 88
getMetricName() (pyspark.ml.evaluation.MulticlassClassificationEvaluator <i>park.ml.classification.NaiveBayes</i> method), 106	getSmoothing() (pyspark.ml.classification.NaiveBayes method), 60
getMetricName() (pyspark.ml.evaluation.RegressionEvaluator <i>park.ml.clustering.PowerIterationClustering</i> method), 105	getSrcCol() (pyspark.ml.clustering.PowerIterationClustering method), 83
getMinDivisibleClusterSize() (pyspark.ml.clustering.BisectingKMeans method), 68	getStagePath() (pyspark.ml.pipeline.PipelineSharedReadWrite static method), 95
getModelType() (pyspark.ml.classification.NaiveBayes method), 60	getStages() (pyspark.ml.pipeline.Pipeline method), 93
getNonnegative() (pyspark.ml.recommendation.ALS method), 87	getStepSize() (pyspark.ml.classification.MultilayerPerceptronClassifier method), 63
getNumFolds() (pyspark.ml.tuning.CrossValidator method), 99	getSubsamplingRate() (pyspark.ml.clustering.LDA method), 77
getNumItemBlocks() (pyspark.ml.classification.LogisticRegression	getThreshold() (pyspark.ml.classification.LogisticRegression

<i>method</i>), 43	I	
getThresholds() (<i>pyspark.ml.classification.LogisticRegression</i> <i>method</i>), 43	intercept (<i>pyspark.ml.classification.LinearSVCModel</i> <i>attribute</i>), 41	
getTopicConcentration() (<i>pyspark.ml.clustering.LDA</i> <i>method</i>), 77	intercept (<i>pyspark.ml.classification.LogisticRegressionModel</i> <i>attribute</i>), 45	
getTopicDistributionCol() (<i>pyspark.ml.clustering.LDA</i> <i>method</i>), 77	intercept (<i>pyspark.ml.regression.AFTSurvivalRegressionModel</i> <i>attribute</i>), 15	
getTrainRatio() (<i>pyspark.ml.tuning.TrainValidationSplit</i> <i>method</i>), 101	intercept (<i>pyspark.ml.regression.GeneralizedLinearRegressionModel</i> <i>attribute</i>), 24	
getUpperBoundsOnCoefficients() (<i>pyspark.ml.classification.LogisticRegression</i> <i>method</i>), 43	intercept (<i>pyspark.ml.regression.LinearRegressionModel</i> <i>attribute</i>), 30	
getUpperBoundsOnIntercepts() (<i>pyspark.ml.classification.LogisticRegression</i> <i>method</i>), 44	interceptVector (<i>pyspark.ml.classification.LogisticRegressionModel</i> <i>attribute</i>), 45	
getUserCol() (<i>pyspark.ml.recommendation.ALS</i> <i>method</i>), 88	isDistributed() (<i>pyspark.ml.clustering.LDAModel</i> <i>method</i>), 80	
getVariancePower() (<i>pyspark.ml.regression.GeneralizedLinearRegression</i> <i>method</i>), 23	isLargerBetter() (<i>pyspark.ml.evaluation.Evaluator</i> <i>method</i>), 103	
H	IsotonicRegression (class in <i>pyspark.ml.regression</i>), 27	
hasSummary (<i>pyspark.ml.classification.LogisticRegressionModel</i> <i>attribute</i>), 45	IsotonicRegressionModel (class in <i>pyspark.ml.regression</i>), 28	
hasSummary (<i>pyspark.ml.clustering.BisectingKMeansModel</i> <i>attribute</i>), 69	itemFactors (<i>pyspark.ml.recommendation.ALSModel</i> <i>attribute</i>), 89	
hasSummary (<i>pyspark.ml.clustering.GaussianMixtureModel</i> <i>attribute</i>), 74	K KMeans (class in <i>pyspark.ml.clustering</i>), 69	
hasSummary (<i>pyspark.ml.clustering.KMeansModel</i> <i>attribute</i>), 72	KMeansModel (class in <i>pyspark.ml.clustering</i>), 71	
hasSummary (<i>pyspark.ml.regression.GeneralizedLinearRegressionModel</i> <i>attribute</i>), 24	KolmogorovSmirnovTest (class in <i>pyspark.ml.stat</i>), 7	
hasSummary (<i>pyspark.ml.regression.LinearRegressionModel</i> <i>attribute</i>), 30	labelCol (<i>pyspark.ml.classification.LogisticRegressionSummary</i> <i>attribute</i>), 46	

attribute), 25
 numInstances (pyspark.ml.regression.LinearRegressionSummary attribute), 32
 numIterations (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 26
 numNonZeros() (pyspark.ml.stat.Summarizer method), 11
O
 objectiveHistory (pyspark.ml.classification.LogisticRegressionTrainingSummary attribute), 48
 objectiveHistory (pyspark.ml.regression.LinearRegressionTrainingSummary attribute), 34
 OneVsRest (class in pyspark.ml.classification), 64
 OneVsRestModel (class in pyspark.ml.classification), 65
P
 ParamGridBuilder (class in pyspark.ml.tuning), 97
 pi (pyspark.ml.classification.NaiveBayesModel attribute), 60
 Pipeline (class in pyspark.ml.pipeline), 93
 PipelineModel (class in pyspark.ml.pipeline), 94
 PipelineModelReader (class in pyspark.ml.pipeline), 94
 PipelineModelWriter (class in pyspark.ml.pipeline), 94
 PipelineReader (class in pyspark.ml.pipeline), 94
 PipelineSharedReadWrite (class in pyspark.ml.pipeline), 94
 PipelineWriter (class in pyspark.ml.pipeline), 95
 PowerIterationClustering (class in pyspark.ml.clustering), 82
 pr (pyspark.ml.classification.BinaryLogisticRegressionSummary attribute), 48
 precisionByLabel (pyspark.ml.classification.LogisticRegressionSummary attribute), 46
 precisionByThreshold (pyspark.ml.classification.BinaryLogisticRegressionSummary attribute), 46
 predict() (pyspark.ml.regression.AFTSurvivalRegressionModel method), 15
 predictionCol (pyspark.ml.classification.LogisticRegressionSummary attribute), 46
 predictionCol (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 predictionCol (pyspark.ml.regression.LinearRegressionSummary attribute), 33
 predictions (pyspark.ml.classification.LogisticRegressionSummary attribute), 47
 predictions (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 predictions (pyspark.ml.regression.IsotonicRegressionModel attribute), 28
 predictions (pyspark.ml.regression.LinearRegressionSummary attribute), 33
 predictQuantiles() (pyspark.ml.regression.AFTSurvivalRegressionModel method), 15
 probability (pyspark.ml.clustering.GaussianMixtureSummary attribute), 75
 probabilityCol (pyspark.ml.classification.LogisticRegressionSummary attribute), 47
 probabilityCol (pyspark.ml.clustering.GaussianMixtureSummary attribute), 75
 pValues (pyspark.ml.regression.GeneralizedLinearRegressionTrainingSummary attribute), 26

pValues (pyspark.ml.regression.LinearRegressionSummary attribute), 32
 pyspark.ml.classification (module), 39
 pyspark.ml.clustering (module), 67
 pyspark.ml.evaluation (module), 103
 pyspark.ml.pipeline (module), 93
 pyspark.ml.recommendation (module), 85
 pyspark.ml.regression (module), 13
 pyspark.ml.stat (module), 5
 pyspark.ml.tuning (module), 97

R

r2 (pyspark.ml.regression.LinearRegressionSummary attribute), 33
 r2adj (pyspark.ml.regression.LinearRegressionSummary attribute), 33
 RandomForestClassificationModel (class in pyspark.ml.classification), 58
 RandomForestClassifier (class in pyspark.ml.classification), 55
 RandomForestRegressionModel (class in pyspark.ml.regression), 36
 RandomForestRegressor (class in pyspark.ml.regression), 34
 rank (pyspark.ml.recommendation.ALSSModel attribute), 89
 rank (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 read() (pyspark.ml.pipeline.Pipeline class method), 93
 read() (pyspark.ml.pipeline.PipelineModel class method), 94
 read() (pyspark.ml.tuning.CrossValidator class method), 99
 read() (pyspark.ml.tuning.CrossValidatorModel class method), 100
 read() (pyspark.ml.tuning.TrainValidationSplit class method), 101
 read() (pyspark.ml.tuning.TrainValidationSplitModel class method), 102
 recallByLabel (pyspark.ml.classification.LogisticRegressionSummary attribute), 47
 recallByThreshold (pyspark.ml.classification.BinaryLogisticRegressionSummary attribute), 49
 recommendForAllItems() (pyspark.ml.recommendation.ALSSModel method), 90
 recommendForAllUsers() (pyspark.ml.recommendation.ALSSModel method), 90
 recommendForItemSubset() (pyspark.ml.recommendation.ALSSModel method), 90
 recommendForUserSubset() (pyspark.ml.recommendation.ALSSModel method), 90
 RegressionEvaluator (class in pyspark.ml.evaluation), 104
 residualDegreeOfFreedom (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 residualDegreeOfFreedomNull (pyspark.ml.regression.GeneralizedLinearRegressionSummary attribute), 25
 residuals (pyspark.ml.regression.LinearRegressionSummary attribute), 33
 residuals() (pyspark.ml.regression.GeneralizedLinearRegressionSummary method), 26
 roc (pyspark.ml.classification.BinaryLogisticRegressionSummary attribute), 49
 rootMeanSquaredError (pyspark.ml.regression.LinearRegressionSummary attribute), 33

S

saveImpl() (pyspark.ml.pipeline.PipelineModelWriter method), 94
 saveImpl() (pyspark.ml.pipeline.PipelineSharedReadWrite

<i>static method</i>), 95		<i>park.ml.regression.IsotonicRegression</i>
<code>saveImpl()</code>	(pys-	<i>method</i>), 27
<i>park.ml.pipeline.PipelineWriter</i>		<code>setFeatureSubsetStrategy()</code>
<i>method</i>), 95		(pys-
<code>scale(pyspark.ml.regression.AFTSurvivalRegressionModel</code>	<i>method</i>), 54	<i>park.ml.classification.GBTClassifier</i>
<i>attribute</i>), 16		<code>setFeatureSubsetStrategy()</code>
<code>scale(pyspark.ml.regression.LinearRegressionModel</code>	<i>park.ml.classification.RandomForestClassifier</i>	(pys-
<i>attribute</i>), 30		<i>method</i>), 57
<code>setAlpha()</code>	(pys-	<code>setFeatureSubsetStrategy()</code>
<i>park.ml.recommendation.ALS</i> <i>method</i>),		(pys-
88		<i>park.ml.regression.GBTRegressor</i>
<code>setBlockSize()</code>	(pys-	<i>method</i>), 19
<i>park.ml.classification.MultilayerPerceptronClassifier</i>	<code>setFeatureSubsetStrategy()</code>	(pys-
<i>method</i>), 63	<i>park.ml.regression.RandomForestRegressor</i>	<i>method</i>), 36
<code>setCensorCol()</code>	(pys-	<code>setFinalStorageLevel()</code>
<i>park.ml.regression.AFTSurvivalRegression</i>		(pys-
<i>method</i>), 15		<i>park.ml.recommendation.ALS</i> <i>method</i>),
		88
<code>setColdStartStrategy()</code>	(pys-	<code>setImplicitPrefs()</code>
<i>park.ml.recommendation.ALS</i> <i>method</i>),		(pys-
88		<i>park.ml.recommendation.ALS</i> <i>method</i>),
		88
<code>setDistanceMeasure()</code>	(pys-	<code>setInitialWeights()</code>
<i>park.ml.clustering.BisectingKMeans</i>		(pys-
<i>method</i>), 68		<i>park.ml.classification.MultilayerPerceptronClassifier</i>
		<i>method</i>), 63
<code>setDistanceMeasure()</code>	(pys-	<code>setInitMode()</code>
<i>park.ml.clustering.KMeans</i> <i>method</i>),		(pys-
71		<i>park.ml.clustering.KMeans</i> <i>method</i>),
		71
<code>setDistanceMeasure()</code>	(pys-	<code>setInitMode()</code>
<i>park.ml.evaluation.ClusteringEvaluator</i>		(pys-
<i>method</i>), 107		<i>park.ml.clustering.PowerIterationClustering</i>
		<i>method</i>), 84
<code>setDocConcentration()</code>	(pys-	<code>setInitSteps()</code>
<i>park.ml.clustering.LDA</i> <i>method</i>),		(pys-
77		<i>park.ml.clustering.KMeans</i> <i>method</i>),
		71
<code>setDstCol()</code>	(pys-	<code>setIntermediateStorageLevel()</code>
<i>park.ml.clustering.PowerIterationClustering</i>		(pyspark.ml.recommendation.ALS
<i>method</i>), 84		<i>method</i>), 88
<code>setEpsilon()</code>	(pys-	<code>setIsotonic()</code>
<i>park.ml.regression.LinearRegression</i>		(pys-
<i>method</i>), 30		<i>park.ml.regression.IsotonicRegression</i>
		<i>method</i>), 28
<code>setFamily()</code>	(pys-	<code>setItemCol()</code>
<i>park.ml.classification.LogisticRegression</i>		(pys-
<i>method</i>), 44		<i>park.ml.recommendation.ALS</i> <i>method</i>),
		88
<code>setFamily()</code>	(pys-	<code>setK()</code>
<i>park.ml.regression.GeneralizedLinearRegression</i>		(pys-
<i>method</i>), 23		<i>park.ml.clustering.BisectingKMeans</i>
		<i>method</i>), 68
<code>setFeatureIndex()</code>	(pys-	<code>setK()</code>
		(pys-

<i>park.ml.clustering.GaussianMixture</i> <i>method</i>), 74	<i>setMetricName()</i> <i>park.ml.evaluation.ClusteringEvaluator</i> <i>method</i>), 107
<i>setK()</i> (<i>pyspark.ml.clustering.KMeans</i> <i>method</i>), 71	<i>setMetricName()</i> (<i>pys-</i> <i>park.ml.evaluation.MulticlassClassificationEvaluator</i> <i>method</i>), 106
<i>setK()</i> (<i>pyspark.ml.clustering.LDA</i> <i>method</i>), 77	<i>setMetricName()</i> (<i>pys-</i> <i>park.ml.evaluation.RegressionEvaluator</i> <i>method</i>), 105
<i>setK()</i> (<i>pys-</i> <i>park.ml.clustering.PowerIterationClustering</i> <i>method</i>), 84	<i>setMinDivisibleClusterSize()</i> (<i>pys-</i> <i>park.ml.clustering.BisectingKMeans</i> <i>method</i>), 69
<i>setKeepLastCheckpoint()</i> (<i>pys-</i> <i>park.ml.clustering.LDA</i> <i>method</i>), 78	<i>setModelType()</i> (<i>pys-</i> <i>park.ml.classification.NaiveBayes</i> <i>method</i>), 60
<i>setLayers()</i> (<i>pys-</i> <i>park.ml.classification.MultilayerPerceptronClassifier</i> <i>method</i>), 63	<i>setNonnegative()</i> (<i>pys-</i> <i>park.ml.recommendation.ALS</i> <i>method</i>), 88
<i>setLearningDecay()</i> (<i>pys-</i> <i>park.ml.clustering.LDA</i> <i>method</i>), 78	<i>setNumBlocks()</i> (<i>pys-</i> <i>park.ml.recommendation.ALS</i> <i>method</i>), 88
<i>setLearningOffset()</i> (<i>pys-</i> <i>park.ml.clustering.LDA</i> <i>method</i>), 78	<i>setNumFolds()</i> (<i>pys-</i> <i>park.ml.tuning.CrossValidator</i> <i>method</i>), 99
<i>setLink()</i> (<i>pys-</i> <i>park.ml.regression.GeneralizedLinearRegression</i> <i>method</i>), 23	<i>setNumItemBlocks()</i> (<i>pys-</i> <i>park.ml.recommendation.ALS</i> <i>method</i>), 89
<i>setLinkPower()</i> (<i>pys-</i> <i>park.ml.regression.GeneralizedLinearRegression</i> <i>method</i>), 23	<i>setNumUserBlocks()</i> (<i>pys-</i> <i>park.ml.recommendation.ALS</i> <i>method</i>), 89
<i>setLinkPredictionCol()</i> (<i>pys-</i> <i>park.ml.regression.GeneralizedLinearRegression</i> <i>method</i>), 23	<i>setOffsetCol()</i> (<i>pys-</i> <i>park.ml.regression.GeneralizedLinearRegression</i> <i>method</i>), 23
<i>setLossType()</i> (<i>pys-</i> <i>park.ml.classification.GBTClassifier</i> <i>method</i>), 55	<i>setOptimizeDocConcentration()</i> (<i>pyspark.ml.clustering.LDA</i> <i>method</i>), 78
<i>setLossType()</i> (<i>pys-</i> <i>park.ml.regression.GBTRegressor</i> <i>method</i>), 19	<i>setOptimizer()</i> (<i>pys-</i> <i>park.ml.clustering.LDA</i> <i>method</i>), 78
<i>setLowerBoundsOnCoefficients()</i> (<i>pyspark.ml.classification.LogisticRegression</i> <i>method</i>), 44	<i>setParams()</i> (<i>pys-</i> <i>park.ml.classification.DecisionTreeClassifier</i> <i>method</i>), 51
<i>setLowerBoundsOnIntercepts()</i> (<i>pys-</i> <i>park.ml.classification.LogisticRegression</i> <i>method</i>), 44	<i>setParams()</i> (<i>pys-</i> <i>park.ml.classification.GBTClassifier</i> <i>method</i>), 55
<i>setMetricName()</i> (<i>pys-</i> <i>park.ml.evaluation.BinaryClassificationEvaluator</i> <i>method</i>), 104	

setParams ()	(pyspark.ml.classification.LinearSVC method), 40	setParams ()	(pyspark.ml.recommendation.ALS method), 89
setParams ()	(pyspark.ml.classification.LogisticRegression method), 44	setParams ()	(pyspark.ml.regression.AFTSurvivalRegression method), 15
setParams ()	(pyspark.ml.classification.MultilayerPerceptronClassifier method), 63	setParams ()	(pyspark.ml.regression.DecisionTreeRegressor method), 17
setParams ()	(pyspark.ml.classification.NaiveBayes method), 60	setParams ()	(pyspark.ml.regression.GBTRegressor method), 19
setParams ()	(pyspark.ml.classification.OneVsRest method), 65	setParams ()	(pyspark.ml.regression.GeneralizedLinearRegression method), 23
setParams ()	(pyspark.ml.classification.RandomForestClassifier method), 58	setParams ()	(pyspark.ml.regression.IsotonicRegression method), 28
setParams ()	(pyspark.ml.clustering.BisectingKMeans method), 69	setParams ()	(pyspark.ml.regression.LinearRegression method), 30
setParams ()	(pyspark.ml.clustering.GaussianMixture method), 74	setParams ()	(pyspark.ml.regression.RandomForestRegressor method), 36
setParams ()	(pyspark.ml.clustering.KMeans method), 71	setParams ()	(pyspark.ml.tuning.CrossValidator method), 99
setParams ()	(pyspark.ml.clustering.LDA method), 79	setParams ()	(pyspark.ml.tuning.TrainValidationSplit method), 101
setParams ()	(pyspark.ml.clustering.PowerIterationClustering method), 84	setQuantileProbabilities ()	(pyspark.ml.regression.AFTSurvivalRegression method), 15
setParams ()	(pyspark.ml.evaluation.BinaryClassificationEvaluator method), 104	setQuantilesCol ()	(pyspark.ml.regression.AFTSurvivalRegression method), 15
setParams ()	(pyspark.ml.evaluation.ClusteringEvaluator method), 107	setRank ()	(pyspark.ml.recommendation.ALS method), 89
setParams ()	(pyspark.ml.evaluation.MulticlassClassificationEvaluator method), 106	setRatingCol ()	(pyspark.ml.recommendation.ALS method), 89
setParams ()	(pyspark.ml.evaluation.RegressionEvaluator method), 105	setSmoothing ()	(pyspark.ml.classification.NaiveBayes method), 40
setParams ()	(pyspark.ml.pipeline.Pipeline method), 111		

`method`), 60
`setSrcCol()` (pyspark.ml.clustering.PowerIterationClustering `method`), 84
`setStages()` (pyspark.ml.pipeline.Pipeline `method`), 93
`setStepSize()` (pyspark.ml.classification.MultilayerPerceptronClassifier `method`), 63
`setSubsamplingRate()` (pyspark.ml.clustering.LDA `method`), 79
`setThreshold()` (pyspark.ml.classification.LogisticRegression `method`), 44
`setThresholds()` (pyspark.ml.classification.LogisticRegression `method`), 44
`setTopicConcentration()` (pyspark.ml.clustering.LDA `method`), 79
`setTopicDistributionCol()` (pyspark.ml.clustering.LDA `method`), 79
`setTrainRatio()` (pyspark.ml.tuning.TrainValidationSplit `method`), 101
`setUpperBoundsOnCoefficients()` (pyspark.ml.classification.LogisticRegression `method`), 44
`setUpperBoundsOnIntercepts()` (pyspark.ml.classification.LogisticRegression `method`), 45
`setUserCol()` (pyspark.ml.recommendation.ALS `method`), 89
`setVariancePower()` (pyspark.ml.regression.GeneralizedLinearRegression `method`), 23
`solver` (pyspark.ml.regression.GeneralizedLinearRegression `attribute`), 26
`subModels` (pyspark.ml.tuning.CrossValidatorModel `attribute`), 100
`subModels` (pyspark.ml.tuning.TrainValidationSplitModel `attribute`), 102
`Summarizer` (class in pyspark.ml.stat), 8
`summary` (pyspark.ml.classification.LogisticRegressionModel `attribute`), 45
`summary` (pyspark.ml.clustering.BisectingKMeansModel `attribute`), 69
`summary` (pyspark.ml.clustering.GaussianMixtureModel `attribute`), 74
`summary` (pyspark.ml.clustering.KMeansModel `attribute`), 72
`summary` (pyspark.ml.regression.GeneralizedLinearRegressionModel `attribute`), 24
`summary` (pyspark.ml.regression.LinearRegressionModel `attribute`), 31
`summary()` (pyspark.ml.stat.SummaryBuilder `method`), 11
`SummaryBuilder` (class in pyspark.ml.stat), 11
T
`test()` (pyspark.ml.stat.ChiSquareTest `static method`), 5
`test()` (pyspark.ml.stat.KolmogorovSmirnovTest `static method`), 7
`theta` (pyspark.ml.classification.NaiveBayesModel `attribute`), 60
`toLocal()` (pyspark.ml.clustering.DistributedLDAModel `method`), 81
`topicsMatrix()` (pyspark.ml.clustering.LDAModel `method`), 80
`TrainingSummary` (pyspark.ml.classification.LogisticRegressionTrainingSummary `attribute`), 48
`totalIterations` (pyspark.ml.classification.LogisticRegressionTrainingSummary `attribute`), 48
`totalIterations` (pyspark.ml.classification.LogisticRegressionTrainingSummary `attribute`), 48

park.ml.regression.LinearRegressionTrainingSummary
attribute), 34

trainingLogLikelihood() (pyspark.ml.clustering.DistributedLDAModel
method), 81

TrainValidationSplit (class in pyspark.ml.tuning), 100

TrainValidationSplitModel (class in pyspark.ml.tuning), 101

trees (pyspark.ml.classification.GBTClassificationModel
attribute), 55

trees (pyspark.ml.classification.RandomForestClassificationModel
attribute), 58

trees (pyspark.ml.regression.GBTRegressionModel
attribute), 20

trees (pyspark.ml.regression.RandomForestRegressionModel
attribute), 37

truePositiveRateByLabel (pyspark.ml.classification.LogisticRegressionSummary
attribute), 47

tValues (pyspark.ml.regression.GeneralizedLinearRegressionTrainingSummary
attribute), 26

tValues (pyspark.ml.regression.LinearRegressionSummary
attribute), 34

U

userFactors (pyspark.ml.recommendation.ALSModel
attribute), 90

V

validateStages() (pyspark.ml.pipeline.PipelineSharedReadWrite
static method), 95

validationMetrics (pyspark.ml.tuning.TrainValidationSplitModel
attribute), 102

variance() (pyspark.ml.stat.Summarizer
static method), 11

vocabSize() (pyspark.ml.clustering.LDAModel
method), 80

weightedFalsePositiveRate (pyspark.ml.classification.LogisticRegressionSummary
attribute), 47

weightedFMeasure() (pyspark.ml.classification.LogisticRegressionSummary
method), 47

weightedPrecision (pyspark.ml.classification.LogisticRegressionSummary
attribute), 47

weightedRecall (pyspark.ml.classification.LogisticRegressionSummary
attribute), 47

weightedTruePositiveRate (pyspark.ml.classification.LogisticRegressionSummary
attribute), 47

weights (pyspark.ml.classification.MultilayerPerceptronClassificationModel
attribute), 63

weights (pyspark.ml.clustering.GaussianMixtureModel
attribute), 74

write() (pyspark.ml.pipeline.Pipeline
method), 94

write() (pyspark.ml.pipeline.PipelineModel
method), 94

write() (pyspark.ml.tuning.CrossValidator
method), 99

write() (pyspark.ml.tuning.CrossValidatorModel
method), 100

write() (pyspark.ml.tuning.TrainValidationSplit
method), 101

write() (pyspark.ml.tuning.TrainValidationSplitModel
method), 102